



HPC2N



Efficient MD simulations at HPC2N

P. Ojeda


Application Expert



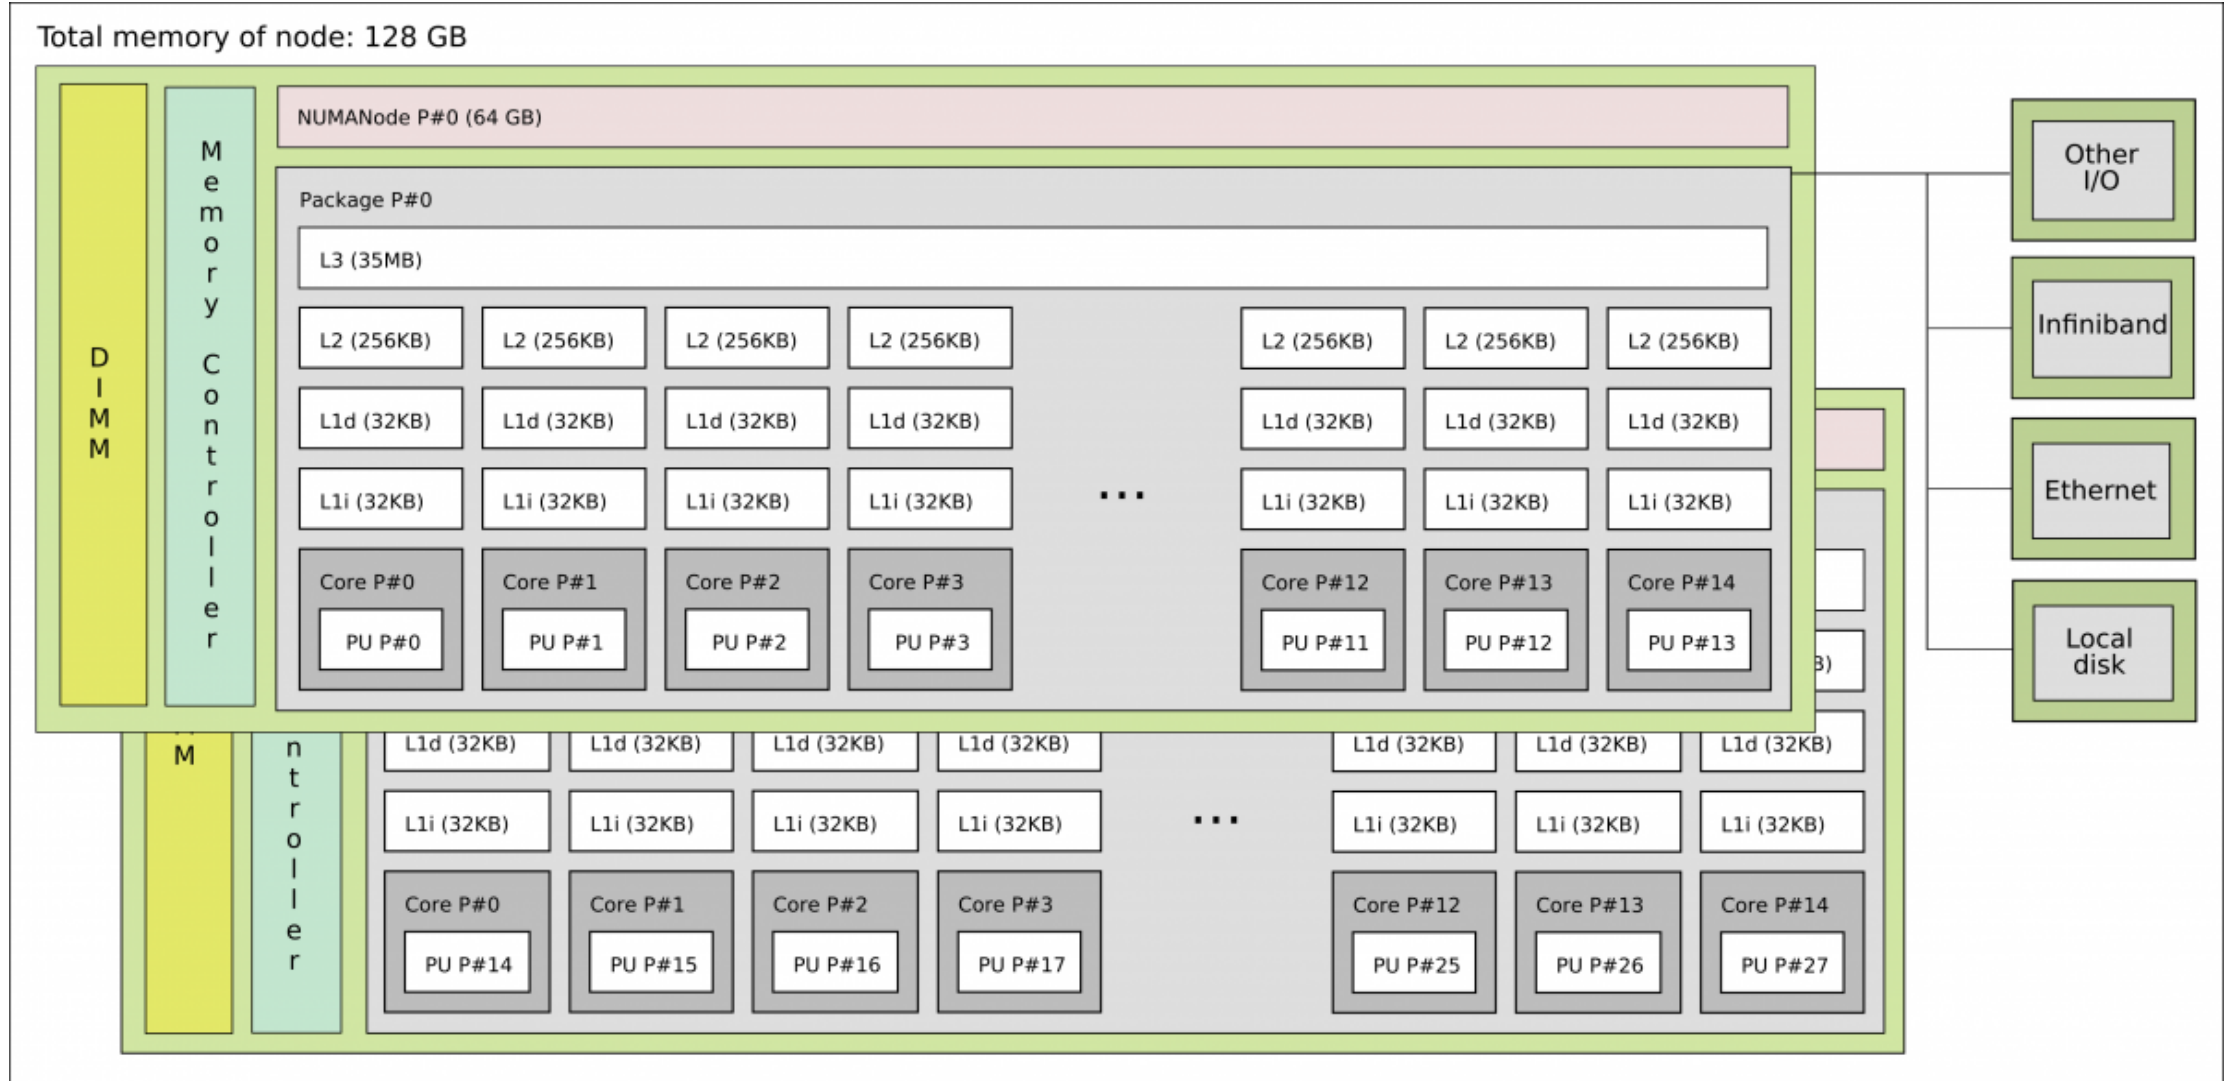
Application Experts at HPC2N:

- Jerry Eriksson
- Åke Sandgren
- Pedro Ojeda (MD, Ab-initio software)

CORE HOURS USED BY DIFFERENT APPLICATIONS (-2021)



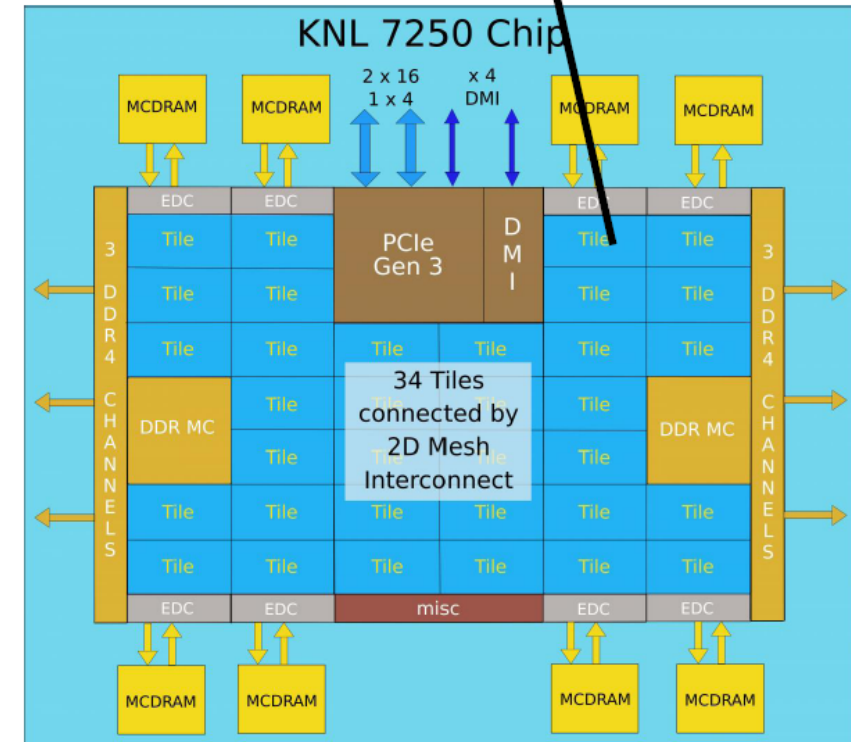

Broadwell node on Kebnekaise



Accelerators



GPU showing the independent units Streaming Multiprocessors (SM).



KNL, composed of several Tiles


$$Z(i) = A * X(i) + Y(i) \quad (\text{Vector Op. SIMD})$$

KNL

Memory Modes


Three Modes. Selected at boot

Cache Mode




- SW-Transparent, Mem-side cache
- Direct mapped. 64B lines.
- Tags part of line
- Covers whole DDR range

Flat Mode



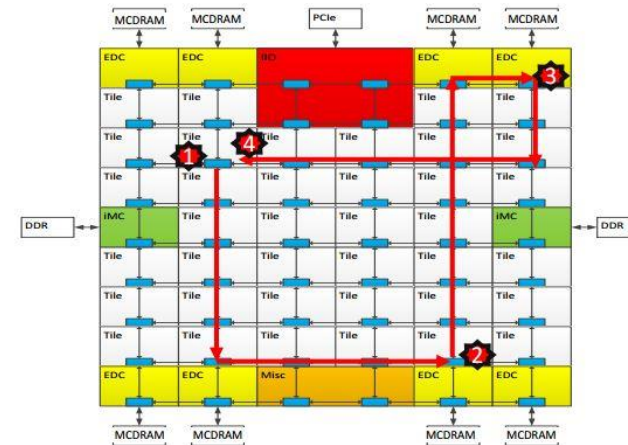
- MCDRAM as regular memory
- SW-Managed
- Same address space

Hybrid Mode



- Part cache, Part memory
- 25% or 50% cache
- Benefits of both

Cluster Mode: All-to-All



Address uniformly hashed across all distributed directories

No affinity between Tile, Directory and Memory

Most general mode. Lower performance than other modes.

Typical Read L2 miss

1. L2 miss encountered
2. Send request to the distributed directory
3. Miss in the directory. Forward to memory
4. Memory sends the data to the requestor

#SBATCH --constraint=a2a,cache

#SBATCH --gres=hbm:4G

KNL

#SBATCH --constraint=quad,flat

numactl -H

node 0 size: 193306 MB

node 0 free: 186258 MB

node 1 cpus:

node 1 size: 16125 MB

node 1 free: 15990 MB

node distances:

node 0 1

0: 10 31

1: 31 10

More information:

<https://www.hpc2n.umu.se/resources/hardware/kebnekaise/knl>

KNL Thread affinity

There are physical 68 cores with 4 hyperthreads on each.



Credits: Intel

Bind the threads by using OpenMP env. var.
export OMP_NUM_THREADS=4
export OMP_PROC_BIND=spread
export OMP_PLACES=cores

srun -n 68 -c 4 --cpu_bind=cores a_knl.out

Alternatively, use Intel var.

KNL Thread affinity

```
#export OMP_PROC_BIND=spread, close, etc.
```

```
#export OMP_PLACES=threads, cores, etc.
```

```
export OMP_NUM_THREADS=4
```

```
srn -n 16 -c 4 --cpu_bind=cores ./xthi
```

```
Hello from rank 0, thread 0, on b-cn1209.hpc2n.umu.se. (core affinity = 0)
```

```
Hello from rank 0, thread 1, on b-cn1209.hpc2n.umu.se. (core affinity = 68)
```

```
Hello from rank 0, thread 2, on b-cn1209.hpc2n.umu.se. (core affinity = 136)
```


```
Hello from rank 0, thread 3, on b-cn1209.hpc2n.umu.se. (core affinity = 204)
```

```
Hello from rank 1, thread 0, on b-cn1209.hpc2n.umu.se. (core affinity = 1)
```

```
Hello from rank 1, thread 1, on b-cn1209.hpc2n.umu.se. (core affinity = 69)
```

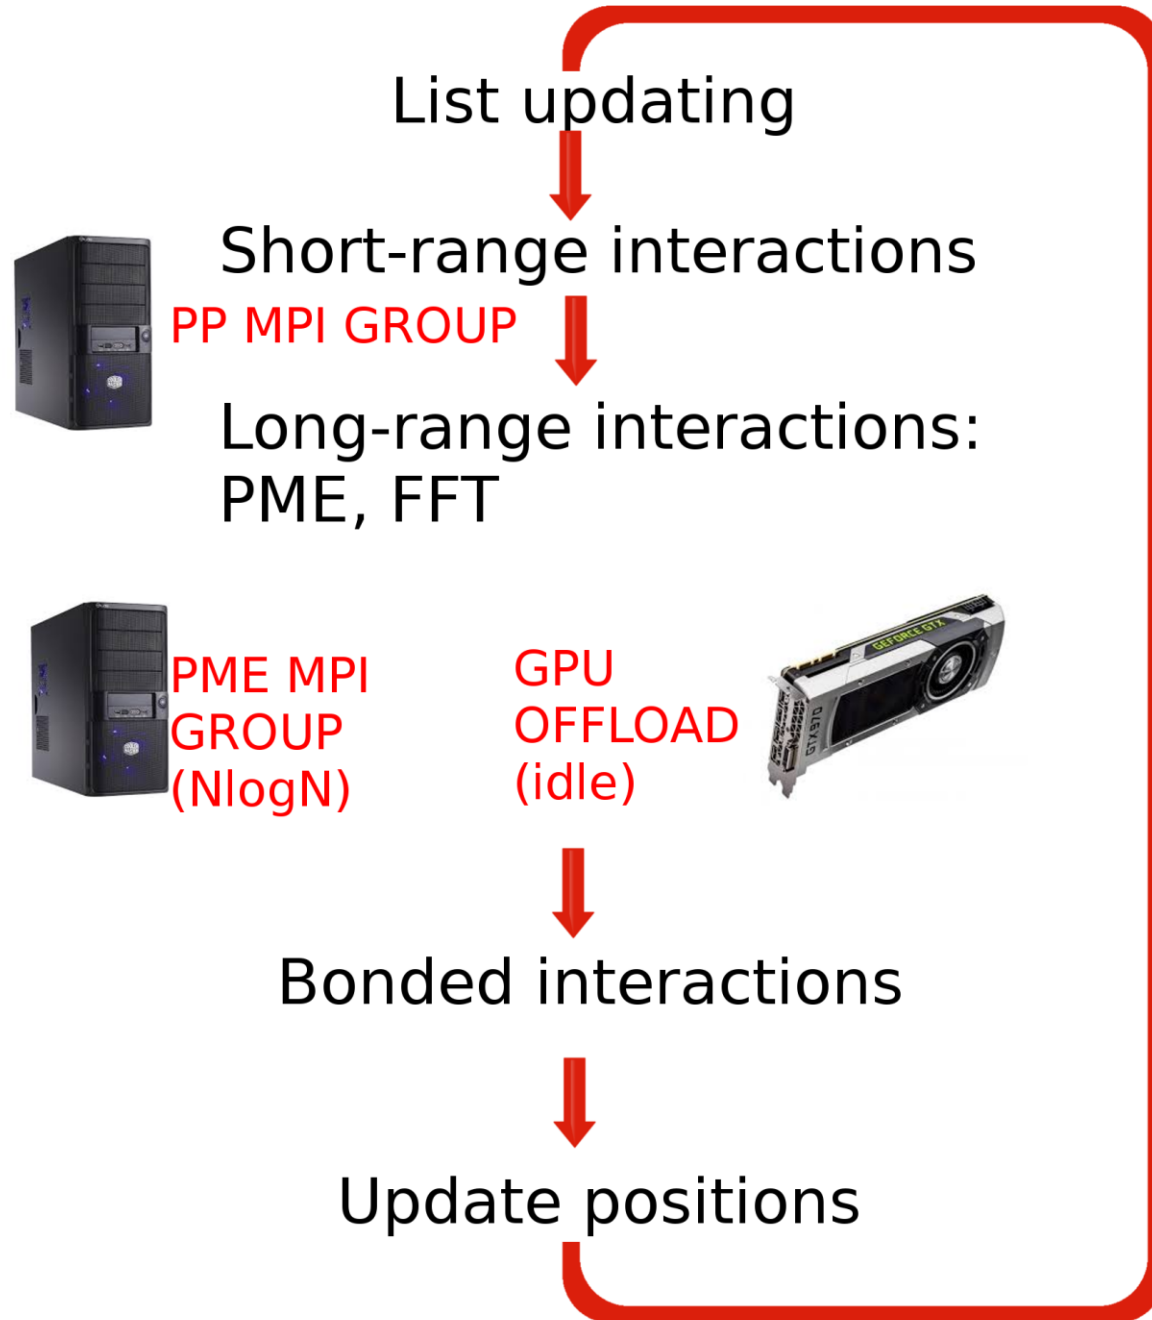
```
Hello from rank 1, thread 2, on b-cn1209.hpc2n.umu.se. (core affinity = 137)
```

```
Hello from rank 1, thread 3, on b-cn1209.hpc2n.umu.se. (core affinity = 205)
```

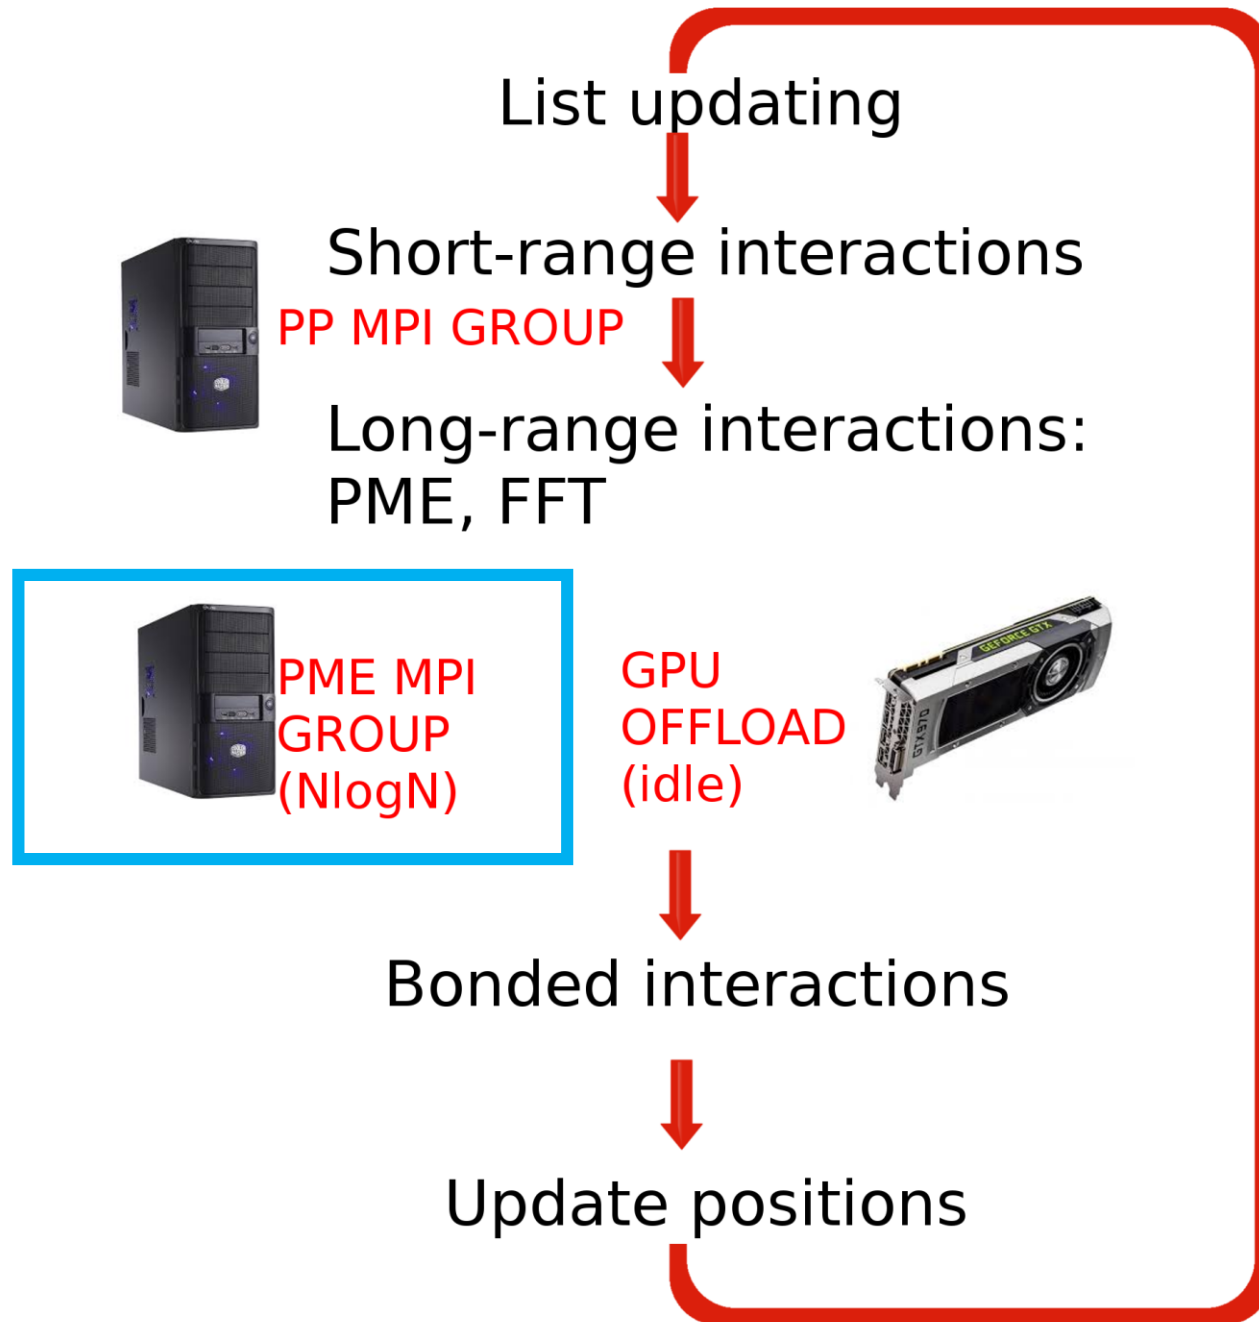


KNL are specially performant if one is developer of the MD application


MD workload



MD workload



Scopus citations

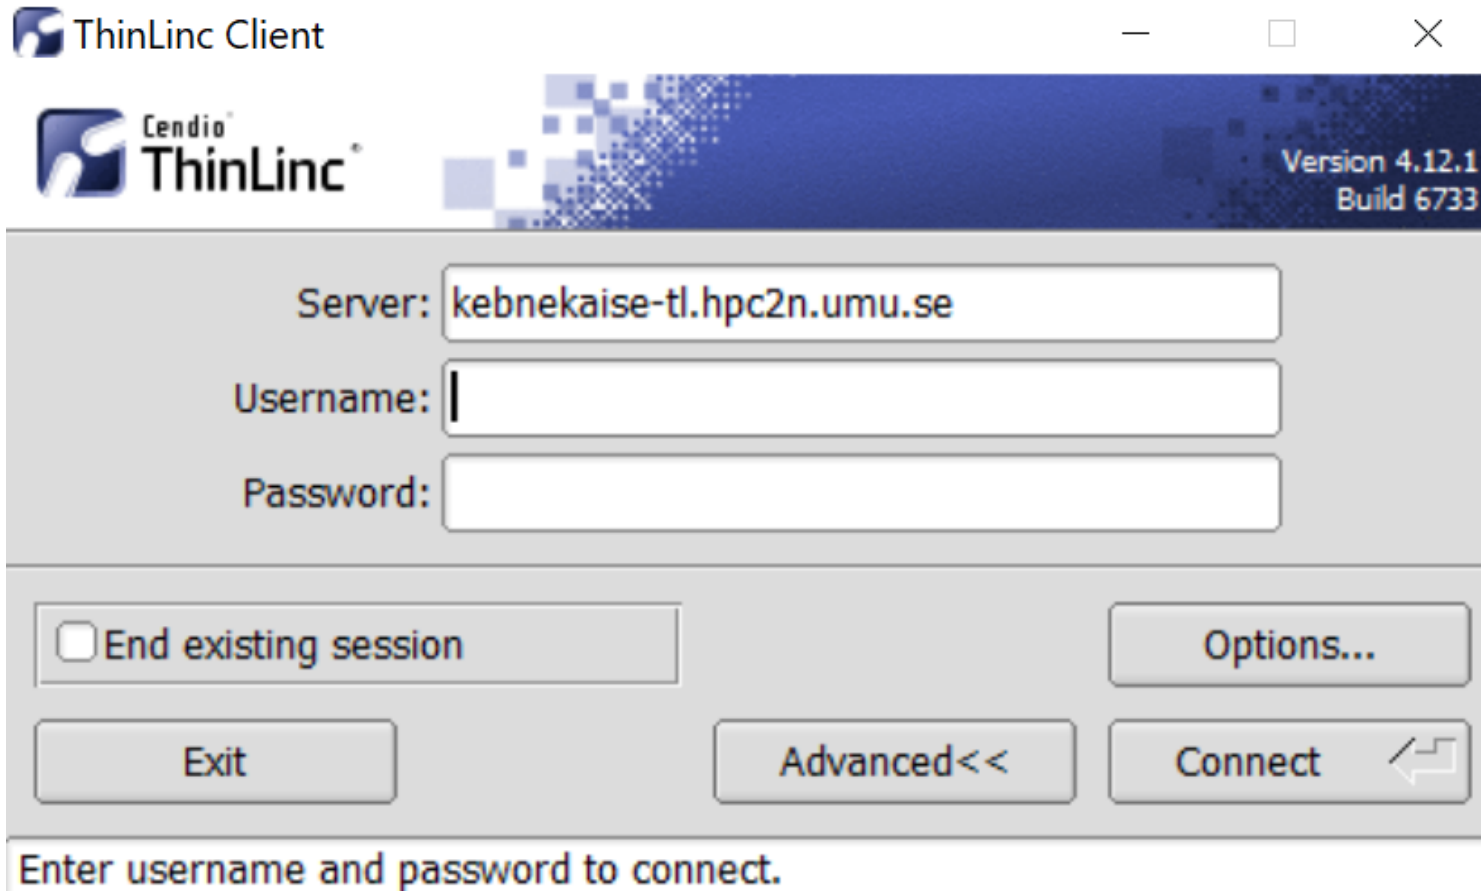


Data from www.scopus.com server by searching for the name of the software. In the case of AMBER the keyword dynamics was added.

ThinLinc client

Download and installation site:

<https://www.cendio.com/thinlinc/download>



The screenshot shows the ThinLinc Client window. The title bar reads "ThinLinc Client". The header area features the Cendio ThinLinc logo on the left and "Version 4.12.1 Build 6733" on the right. The main interface has three input fields: "Server:" with the text "kebnekaise-tl.hpc2n.umu.se", "Username:" which is empty, and "Password:" which is empty. Below these fields are four buttons: "End existing session" (with an unchecked checkbox), "Options...", "Exit", and "Advanced<<". To the right of the "Advanced<<" button is a "Connect" button with a right-pointing arrow icon. At the bottom of the window, a status bar contains the text "Enter username and password to connect."

Username and Password are the ones you received for HPC2N

Just the first time you use ThinLinc, got to Options -> Screen and uncheck the box "Full screen mode"



Benchmark

- Solvated protein (1AKI)
- 158945 atoms
- 1.2 nm cutoff radius
- 1 fs time step
- PME for electrostatics
- Input scripts generated in <http://www.charmm-gui.org/>

AMBER

- Collection of independent routines
- It uses Sander/PMEMD for solving the Newton's equations
- It offers a robust set of analysis tools

Amber12 throughput JAC NVE Benchmark



```
ml GCC/9.3.0 OpenMPI/4.0.3
ml Amber/18.17-AmberTools-19.12-Python-2.7.18
srun pmemd.MPI -O -i input.mdin
srun pmemd.cuda.MPI -O -i input.mdin
```


AMBER Tools

- For setting up a simulation (initial structure, solvation, ions, ...):

Load the modules:

```
$ml GCC/9.3.0 OpenMPI/4.0.3
```

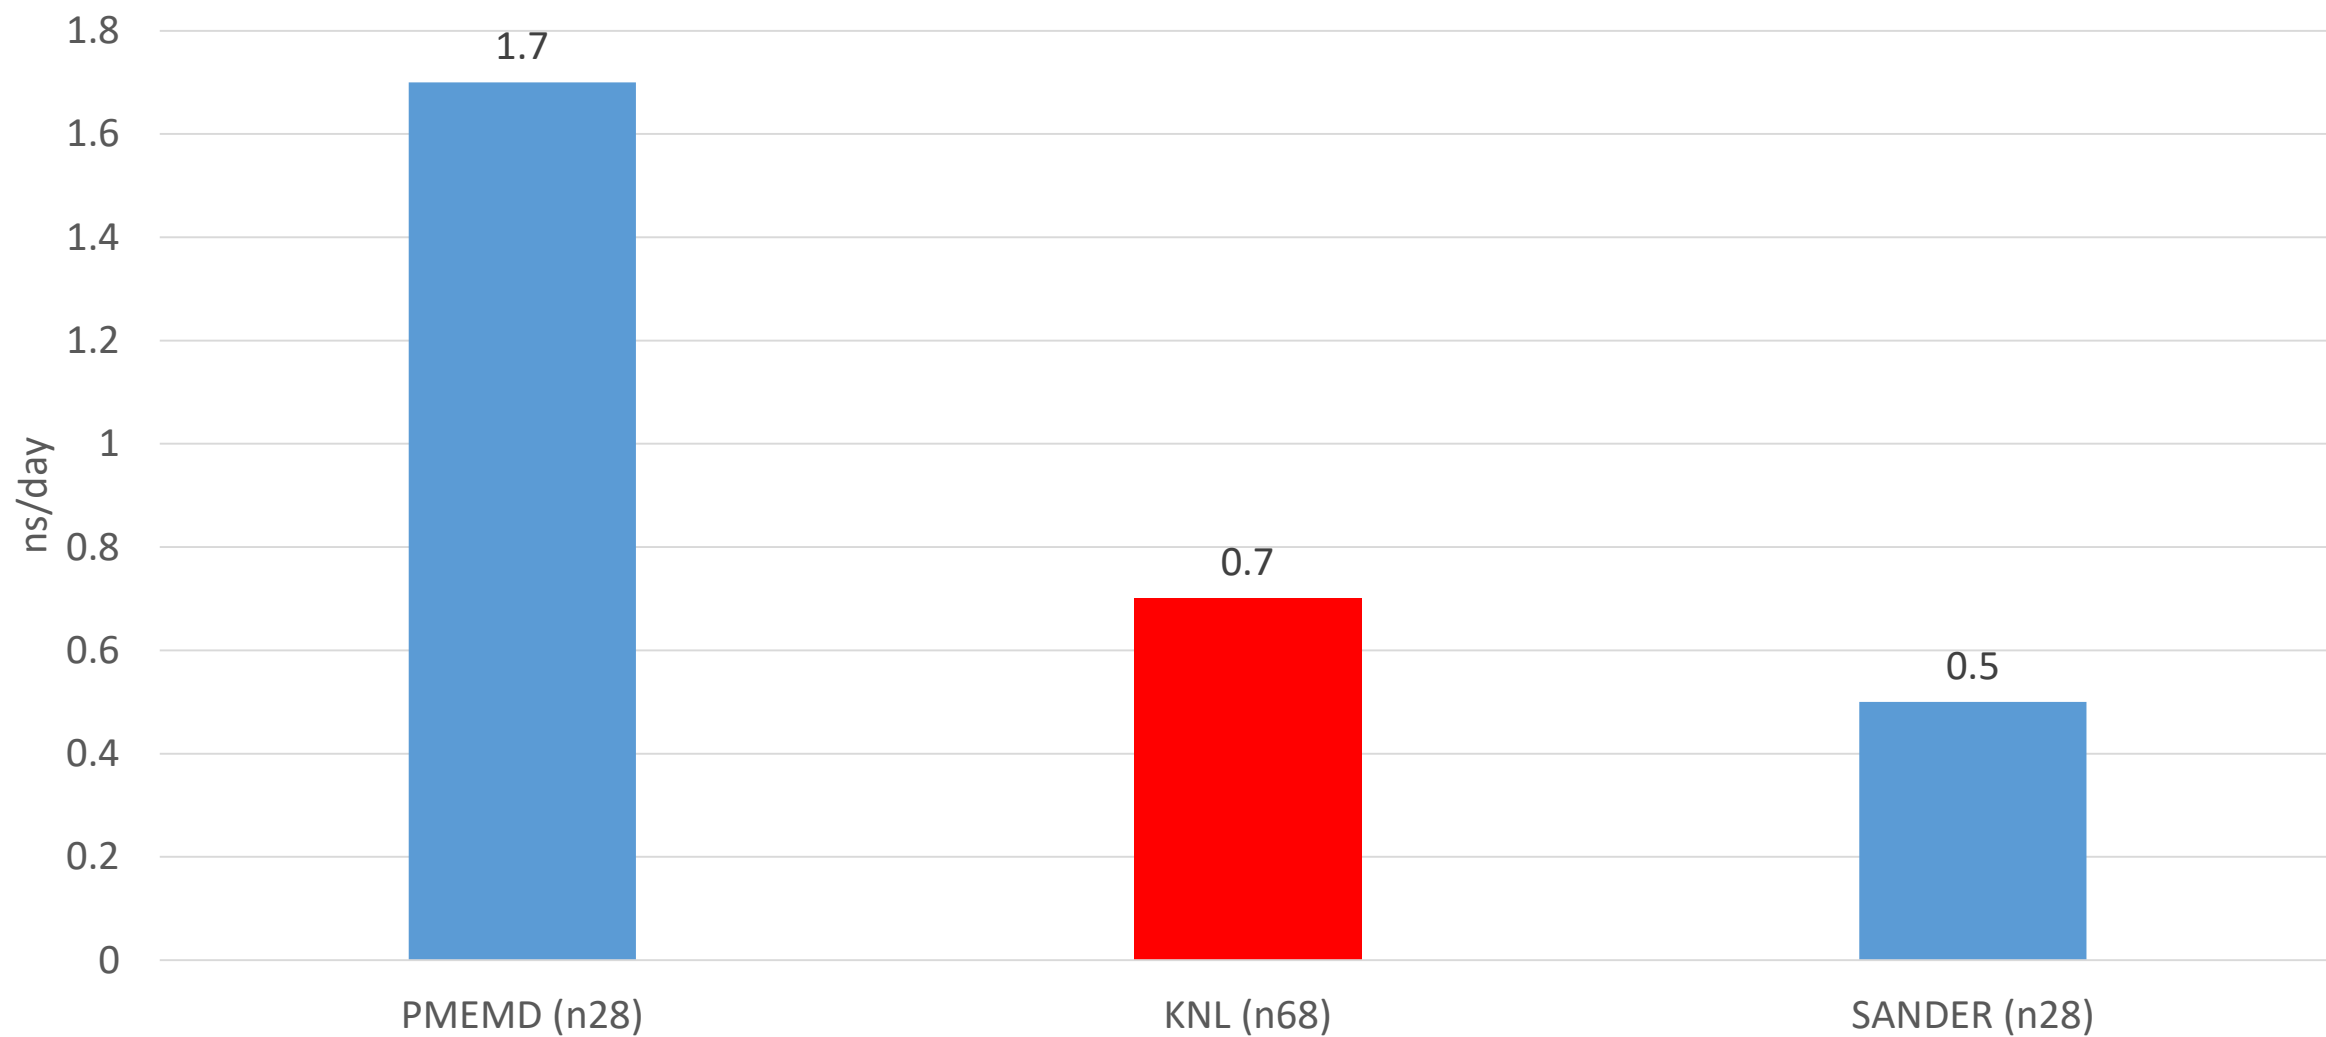
```
$ml Amber/18.17-AmberTools-19.12-Python-2.7.18
```

On the command line:

```
> tleap, antechamber, cpptraj, ...
```

AMBER

AMBER



AMBER-GPU

```
#!/bin/bash
#SBATCH -A staff
#SBATCH -t 00:50:00
#SBATCH -N 1
#SBATCH -n 4
#SBATCH --gres=gpu:k80:2
#SBATCH -p batch
#SBATCH --exclusive
#SBATCH --output=job_str.out
#SBATCH --error=job_str.err
#SBATCH --mail-type=END
```

```
ml purge > /dev/null 2>&1
```

```
ml GCC/9.3.0 OpenMPI/4.0.3
```


```
ml Amber/18.17-AmberTools-19.12-Python-2.7.18
```

```
srun pmemd.cuda.MPI -O -i input.mdin -p input.parm7 -c input.rst7 -o input.mdout
```

AMBER

On Kebnekaise, best performance is achieved with 4 MPIs/Node and using 2 GPU cards

Nr. Ranks (-n)	Nr. GPUs	ns/day
2	1	6
2	2	6
4	2	8
8 (single-node)	2	6
8 (multi-node)	2	10



AMBER

- If you observe any issue with a GPU run, go back to the pure CPU version (AMBER 2018 user guide)
- In case you want to perform independent simulations use the variable `CUDA_VISIBLE_DEVICES` to specify the GPU you will use or better request a single GPU card (`--gres=gpu:k80:1`)
- GPU enhanced sampling: Gaussian Accelerated MD (JCTC, 11, 3584-3595, 2015)

AMBER

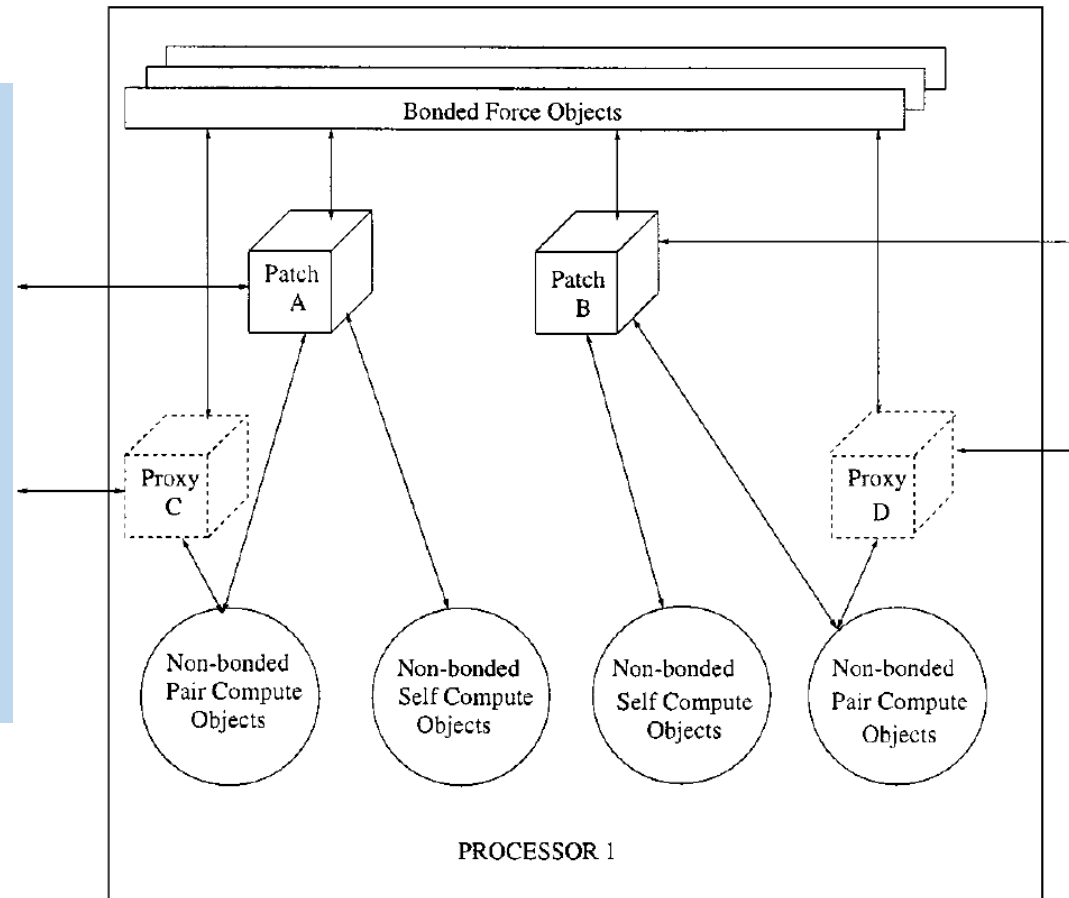
Resources

AMBER tutorials

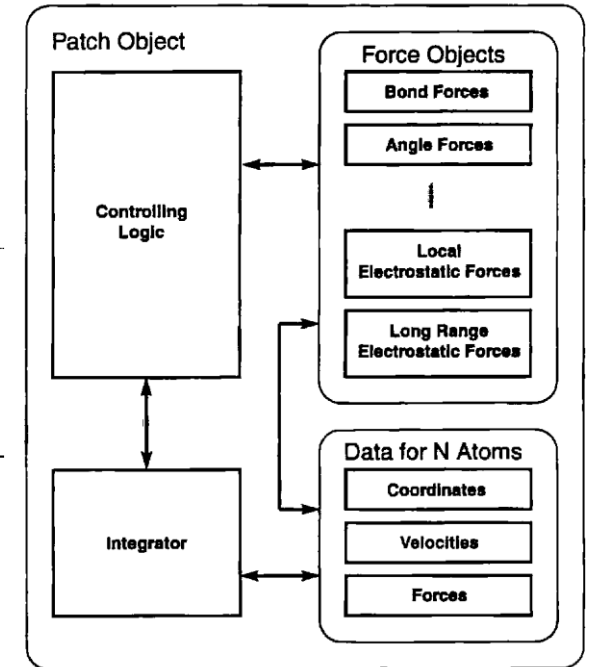
<http://ambermd.org/tutorials/>

NAMD

- Based on charm++ communication protocol
- It is object-oriented
- Versions: **single node, multi-node, GPU, and KNL**
- Highly scalable
- Message driven comm.



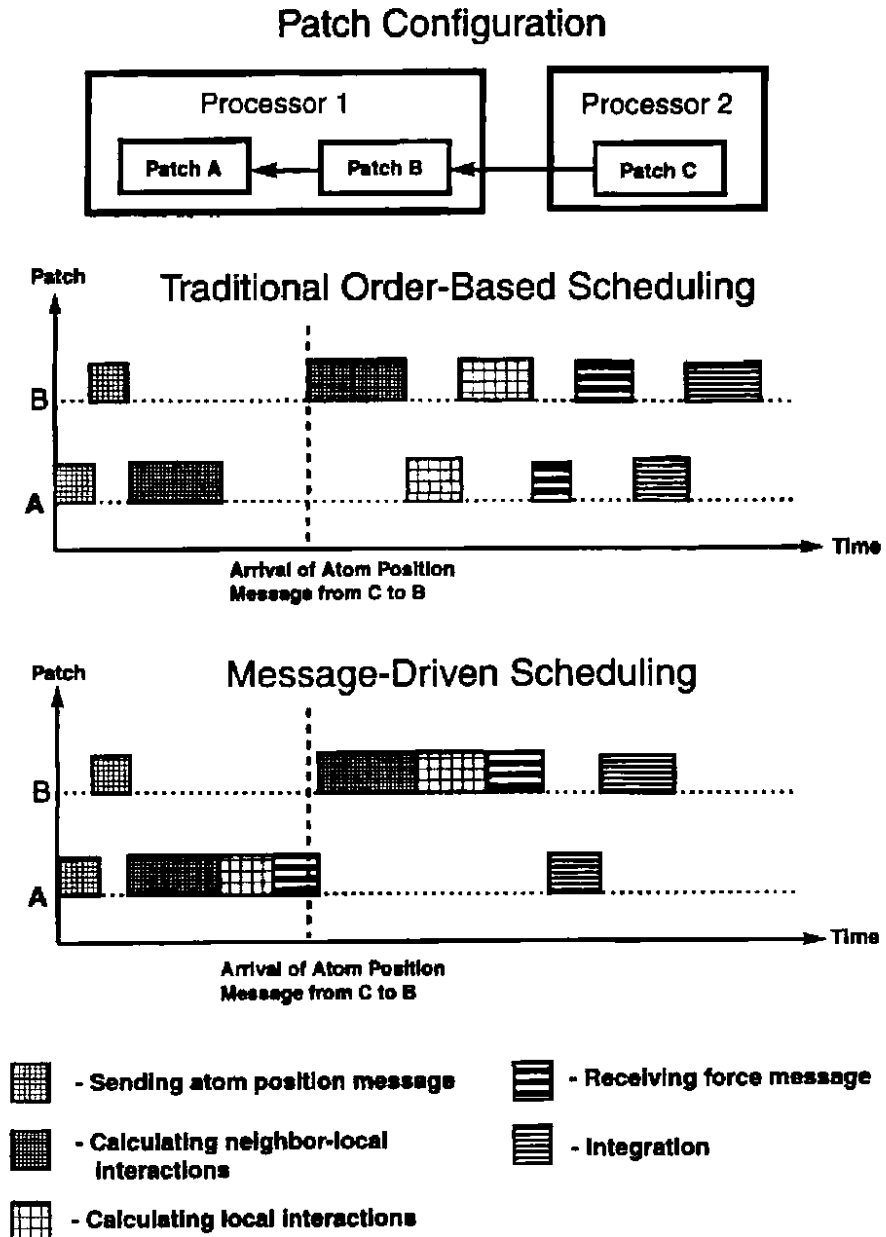
Credits: JCC, 151, 283 (1999)



Credits: The Int. J. Supercomp. Appl. And High Performance Comp., 10, 251-268 (1996)

NAMD

- Based on charm++ communication protocol
- It is object-oriented
- Versions: **single node, multi-node, GPU, and KNL**
- Highly scalable
- Message driven comm.



NAMD (multiple nodes)

```
#!/bin/bash
#SBATCH -A SNICyyyyy-xx-zz
#Asking for 10 min.
#SBATCH -t 00:10:00
#Number of nodes
#SBATCH -N 2
#Ask for 56 processes (2x28 cores on the nodes)
#SBATCH -n 56
#SBATCH --exclusive
#Load modules necessary for running NAMD
ml GCC/9.3.0 OpenMPI/4.0.3
ml NAMD/2.14-mpi
srun namd2 config_file > output_file
```

NAMD (GPU) (single node)

```
#!/bin/bash
#SBATCH -A SNICyyyyy-xx-zz
#SBATCH -t 00:50:00
#SBATCH -N 1
#SBATCH -n 28
#SBATCH --exclusive
#Ask for 2 GPU cards
#SBATCH --gres=gpu:k80:2
#Load modules necessary for running NAMD
ml GCC/9.3.0 CUDA/11.0.2 OpenMPI/4.0.3
ml NAMD/2.14-nompi
#Execute NAMD
namd2 +p 28 +setcpuaffinity +idlepoll +devices $CUDA_VISIBLE_DEVICES config > output.dat
```

NAMD

- NAMD will scale if the number of patches >> number of processes

```
Info: Startup phase 5 took 0.00011301 s, 1131.09 MB of memory in use  
Info: PATCH GRID IS 7 (PERIODIC) BY 7 (PERIODIC) BY 7 (PERIODIC)
```

- One can dedicate some processors to solve PME long-range part (config file)

```
PMEProcessors 8  
Ldb unload PME yes
```

- For the CUDA version, one can offload PME to GPUs:

```
usePMECUDA on  
PMEoffload on
```

```
In the output file you will see:  
Info: PME RECIPROCAL SUM OFFLOADED TO GPU
```

- Colvars module for free energy calculations can be run on GPUs.

More information:

<https://www.ks.uiuc.edu/Research/namd/2.13/ug/node106.html>

<https://www.ks.uiuc.edu/Research/namd/wiki/?NamdPerformanceTuning>

NAMD


- Colvars module for free energy calculations can be run on GPUs

```
colvar {  
  name restrain01  
  width 0.5  
  lowerboundary 0.0  
  upperboundary 8.0  
  lowerwallconstant 100.0  
  upperwallconstant 100.0  
  
  distanceXY {  
    main { atomnumbers {list of atoms} }  
    ref { atomnumbers {list of atoms} }  
    axis (0.0, 0.0, 1.0)  
  }  
}
```

Restraint where a channel
is asymmetric

NAMD

Use the **+setcpuaffinity**
flag for a **10% speedup**
for the CUDA version




NAMD

With the new Volta GPU cards one
can speed up the simulation by
1.6x

#SBATCH --gres=gpu:v100:2

SMP (Skylake)	
Setup	Timing (ns/day)
+p28 (2GPU)	15.1
+p14 (1GPU)	9.8

MTS=Multiple Time Step



NAMD

Resources

NAMD tutorials:

<http://www.ks.uiuc.edu/Training/Tutorials/#namd>

GROMACS

```
#!/bin/bash
#SBATCH -A SNICyyyyy-xx-zz
#SBATCH -t 00:10:00
#SBATCH -n 4
#SBATCH -c 7
# Asking for 2 GPUs
#SBATCH --gres=gpu:k80:2
```

GROMACS recognizes the number of available GPU cards

```
ml GCC/5.4.0-2.26 OpenMPI/2.0.1 CUDA/8.0.44
ml GROMACS/2016-hybrid
```

```
if [ -n "$SLURM_CPUS_PER_TASK" ]; then
    mdargs="-ntomp $SLURM_CPUS_PER_TASK"
else
    mdargs="-ntomp 1"
fi
```

```
srun gmx_mpi mdrun $mdargs -npme 0 -dlb yes -v -deffnm step4.1_equilibration
```


gmx tune_pme

Individual timings for input file 0 (npt_bench00.tpr):

PME ranks	Gcycles	ns/day	PME/f	Remark
0	4355.019	57.776	-	OK.
0	4547.105	55.335	-	OK.
0	4289.420	58.659	-	OK.
-1(0)	4455.791	56.469	-	OK.
-1(0)	4440.157	56.668	-	OK.
-1(0)	4275.551	58.850	-	OK.

Tuning took 7.7 minutes.

Summary of successful runs:

Line	tpr	PME ranks	Gcycles	Av.	Std.dev.	ns/day	PME/f	DD grid
0	0	0	4397.181	133.917	57.257	-	4 1 1	
1	0	-1(0)	4390.500	99.855	57.329	-	4 1 1	

Best performance was achieved with the automatic number of PME ranks (see line 1)

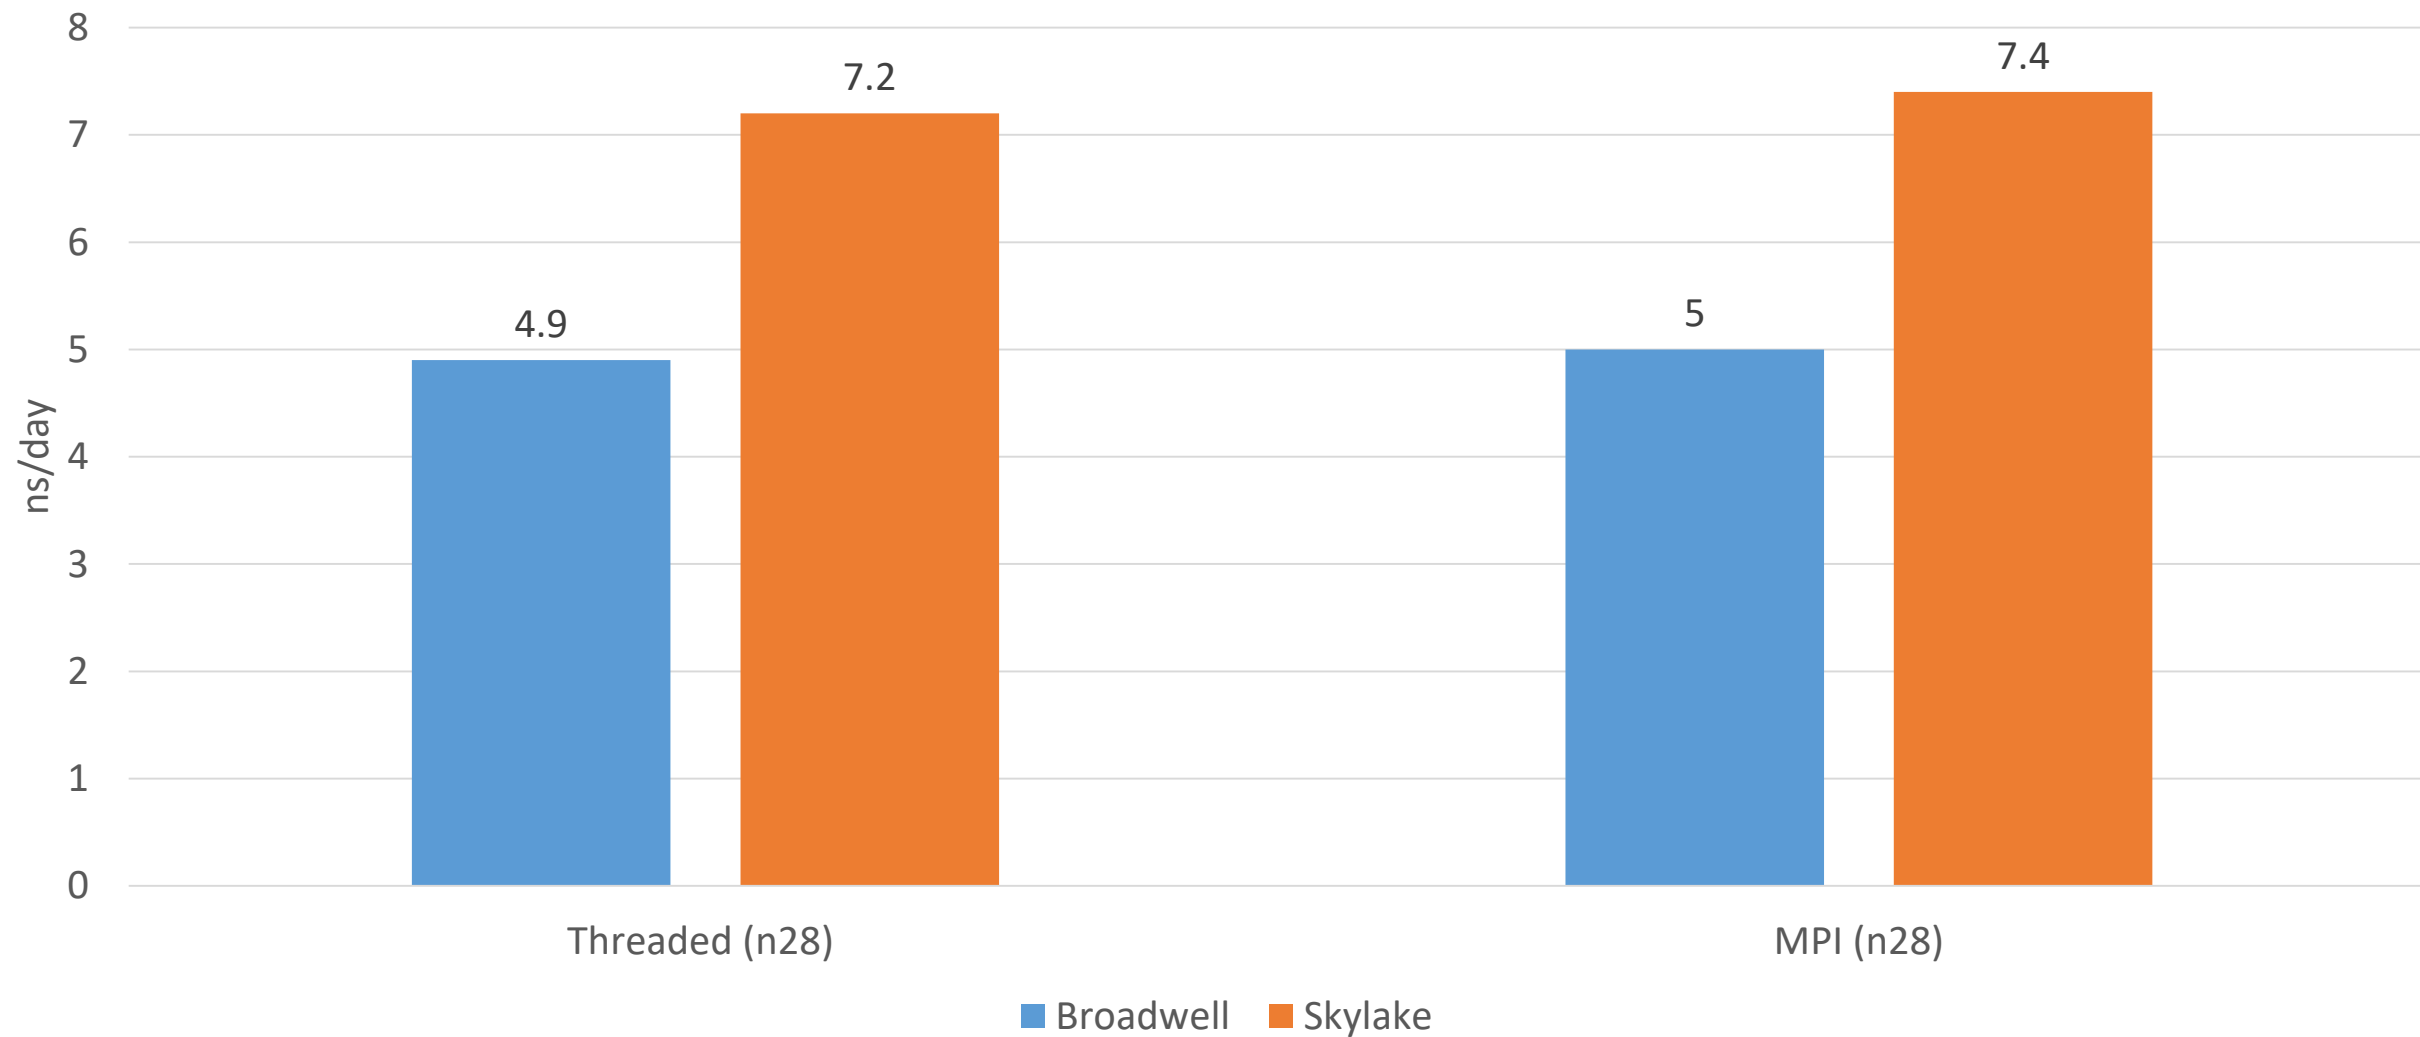
Please use this command line to launch the simulation:

```
mpirun -np 4 gmx_mpi mdrun -npme -1 -s npt.tpr -ntomp 7 -dlb yes
```

GROMACS

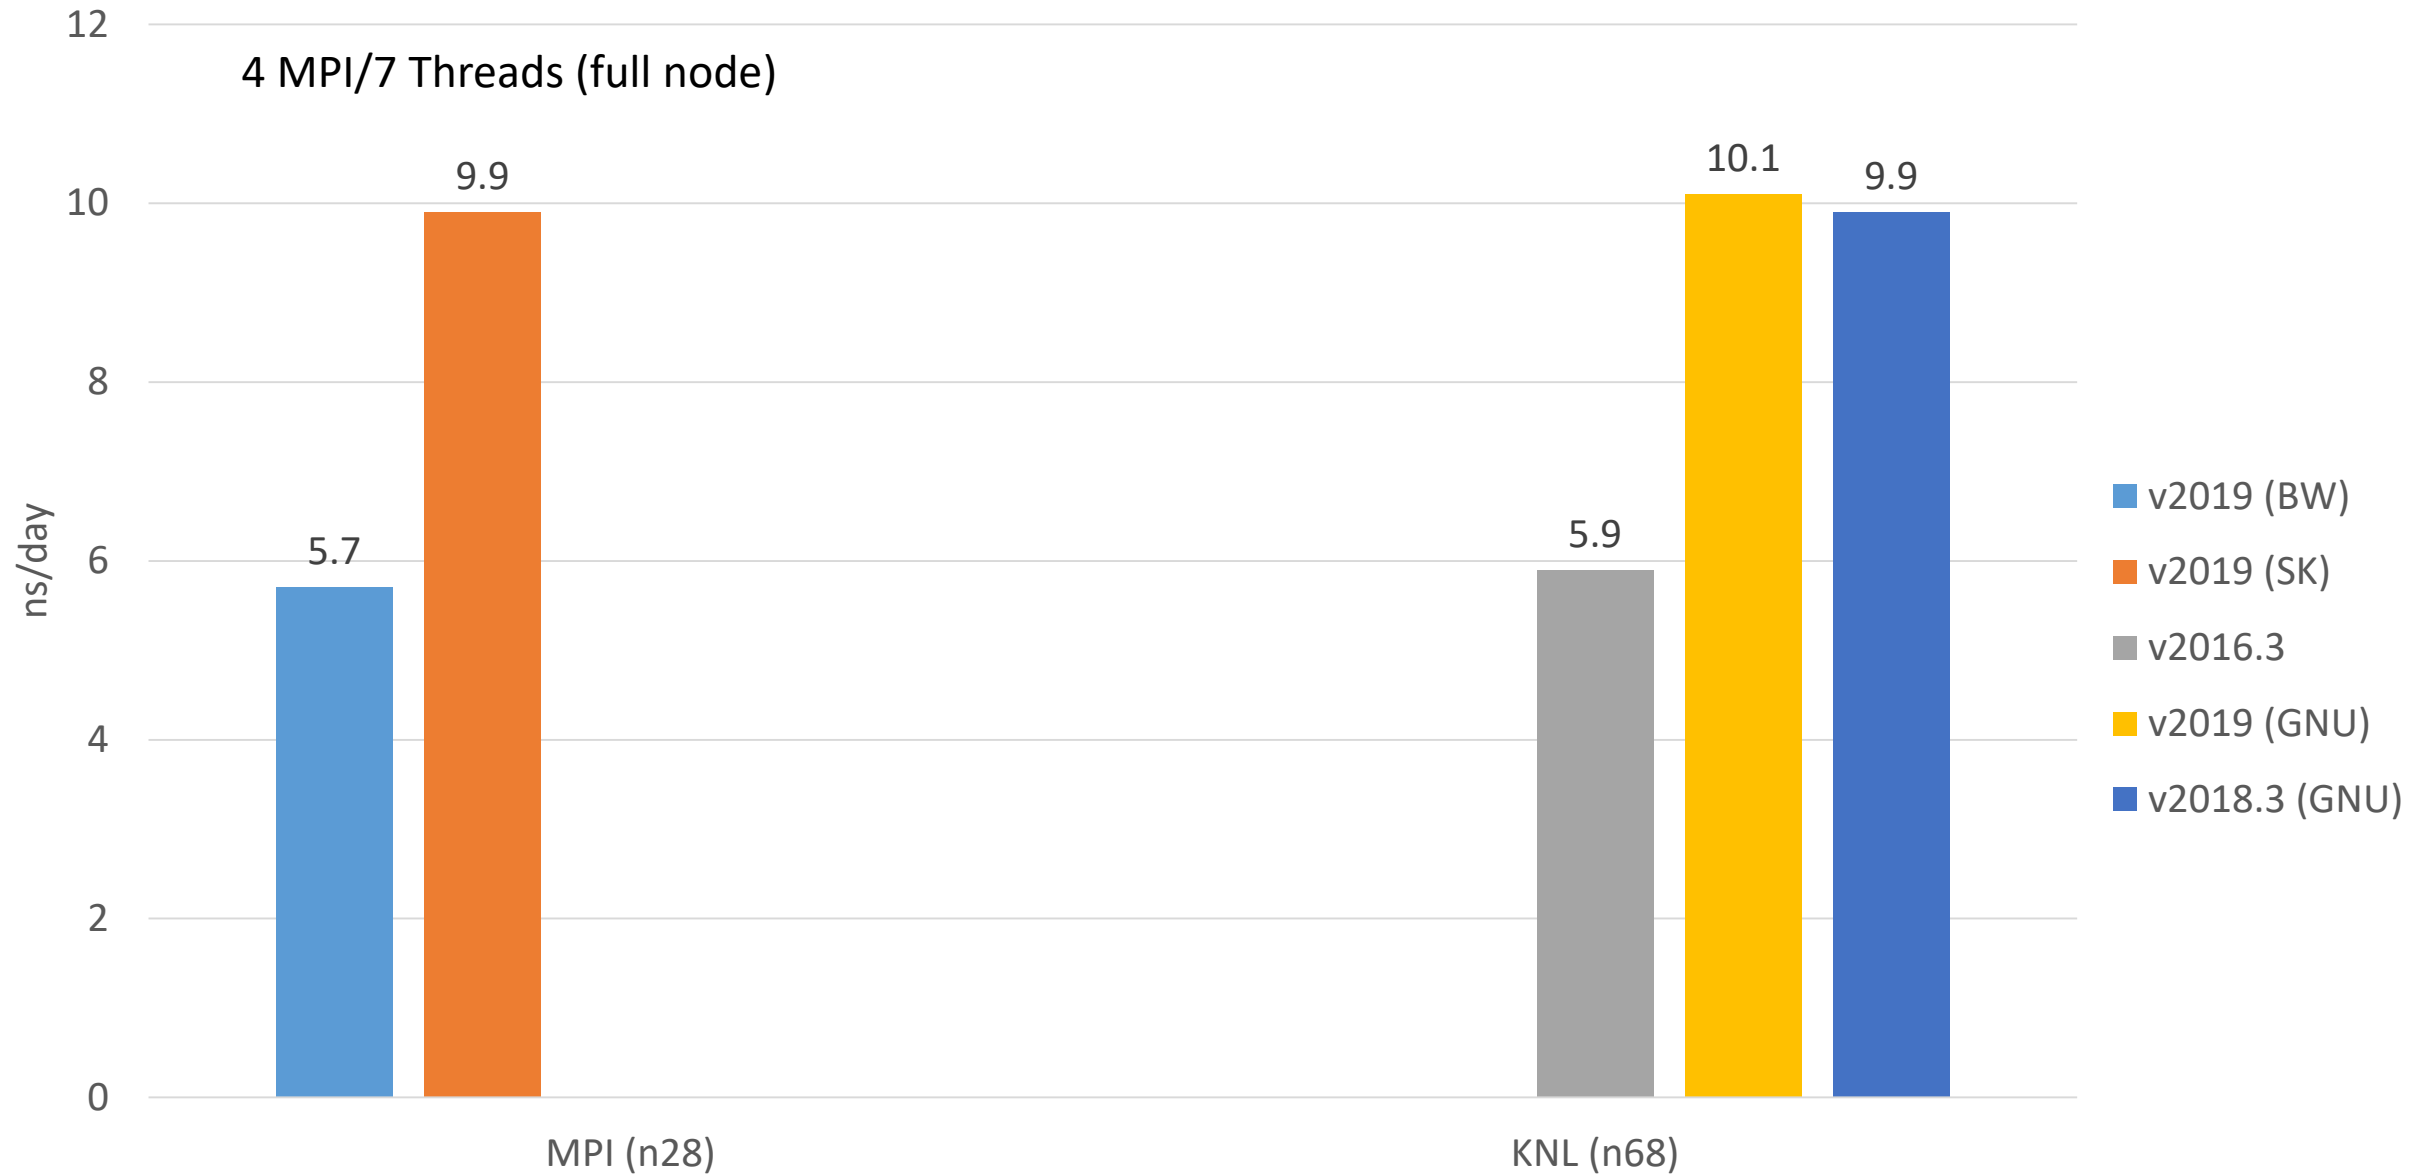
4 MPI/7 Threads (full node)

GROMACS/2016.3



GROMACS

GROMACS BW/SK vs. KNL



GROMACS KNL

```
#!/bin/bash
```

```
#SBATCH -A project_ID
```

```
#SBATCH -t 00:50:00
```

```
#SBATCH -N 1
```

```
#SBATCH -n 68
```

```
#SBATCH -p knl
```

```
#SBATCH --constraint=cache,quad
```

```
#SBATCH --exclusive
```

```
ml GCC/7.3.0-2.30 OpenMPI/3.1.1
```

```
ml GROMACS/2018.3
```

```
export OMP_NUM_THREADS=2
```

```
gmx mdrun -ntmpi 68 -npme 18 -ntomp 2 -pin on -pinoffset 0 -pinstride 2 -dlb auto -v -deffnm step4.1_eq
```


GROMACS KNL

- Login to kebnekaise-knl.hpc2n.umu.se to check the GROMACS modules available and their dependencies.
- Submit your script either from the KNL or from the standard login nodes
- KNL queue is most of the time available compared to GPU one

GROMACS KNL

GROMACS/2019 (Threaded-MPI)


2,136,412 atoms
ts = 4fs
It uses virtual sites
cutoff radii = 1.0nm



Benchmark (2M atoms) from: <https://www.mpibpc.mpg.de/grubmueller/bench>

GROMACS KNL


GROMACS/2019 (Intel, multi-node)



Benchmark (2M atoms) from: <https://www.mpibpc.mpg.de/grubmueller/bench>

GROMACS (toolchain effects)

GROMACS/2019 –KNL (Threaded-MPI)



GROMACS (single node, K80)

- (-n2,-c14) `srun gmx_mpi mdrun -ntomp 14 -dlb yes` will use 2 GPUs engines

Using 2 MPI processes

Using 14 OpenMP threads per MPI process

On host b-cn1309.hpc2n.umu.se 2 GPUs auto-selected for this run.

Mapping of GPU IDs to the 2 GPU tasks in the 2 ranks on this node:

PP:0,PP:1

- (-n4,-c7) `srun gmx_mpi mdrun -ntomp 7 -dlb yes` will use 4 GPUs engines

Using 4 MPI processes

Using 7 OpenMP threads per MPI process

On host b-cn1309.hpc2n.umu.se 4 GPUs auto-selected for this run.

Mapping of GPU IDs to the 4 GPU tasks in the 4 ranks on this node:

PP:0,PP:1,PP:2,PP:3

GROMACS (single node, K80)

- (-n4,-c7) `gmx mdrun -ntmpi 4 -ntomp 7 -dlb yes` will use 4 GPUs engines

Using 4 MPI threads

Using 7 OpenMP threads per tMPI thread

On host b-cn1105.hpc2n.umu.se 4 GPUs auto-selected for this run.


Mapping of GPU IDs to the 4 GPU tasks in the 4 ranks on this node:

PP:0,PP:1,PP:2,PP:3

PP tasks will do (non-perturbed) short-ranged and most bonded interactions on the GPU

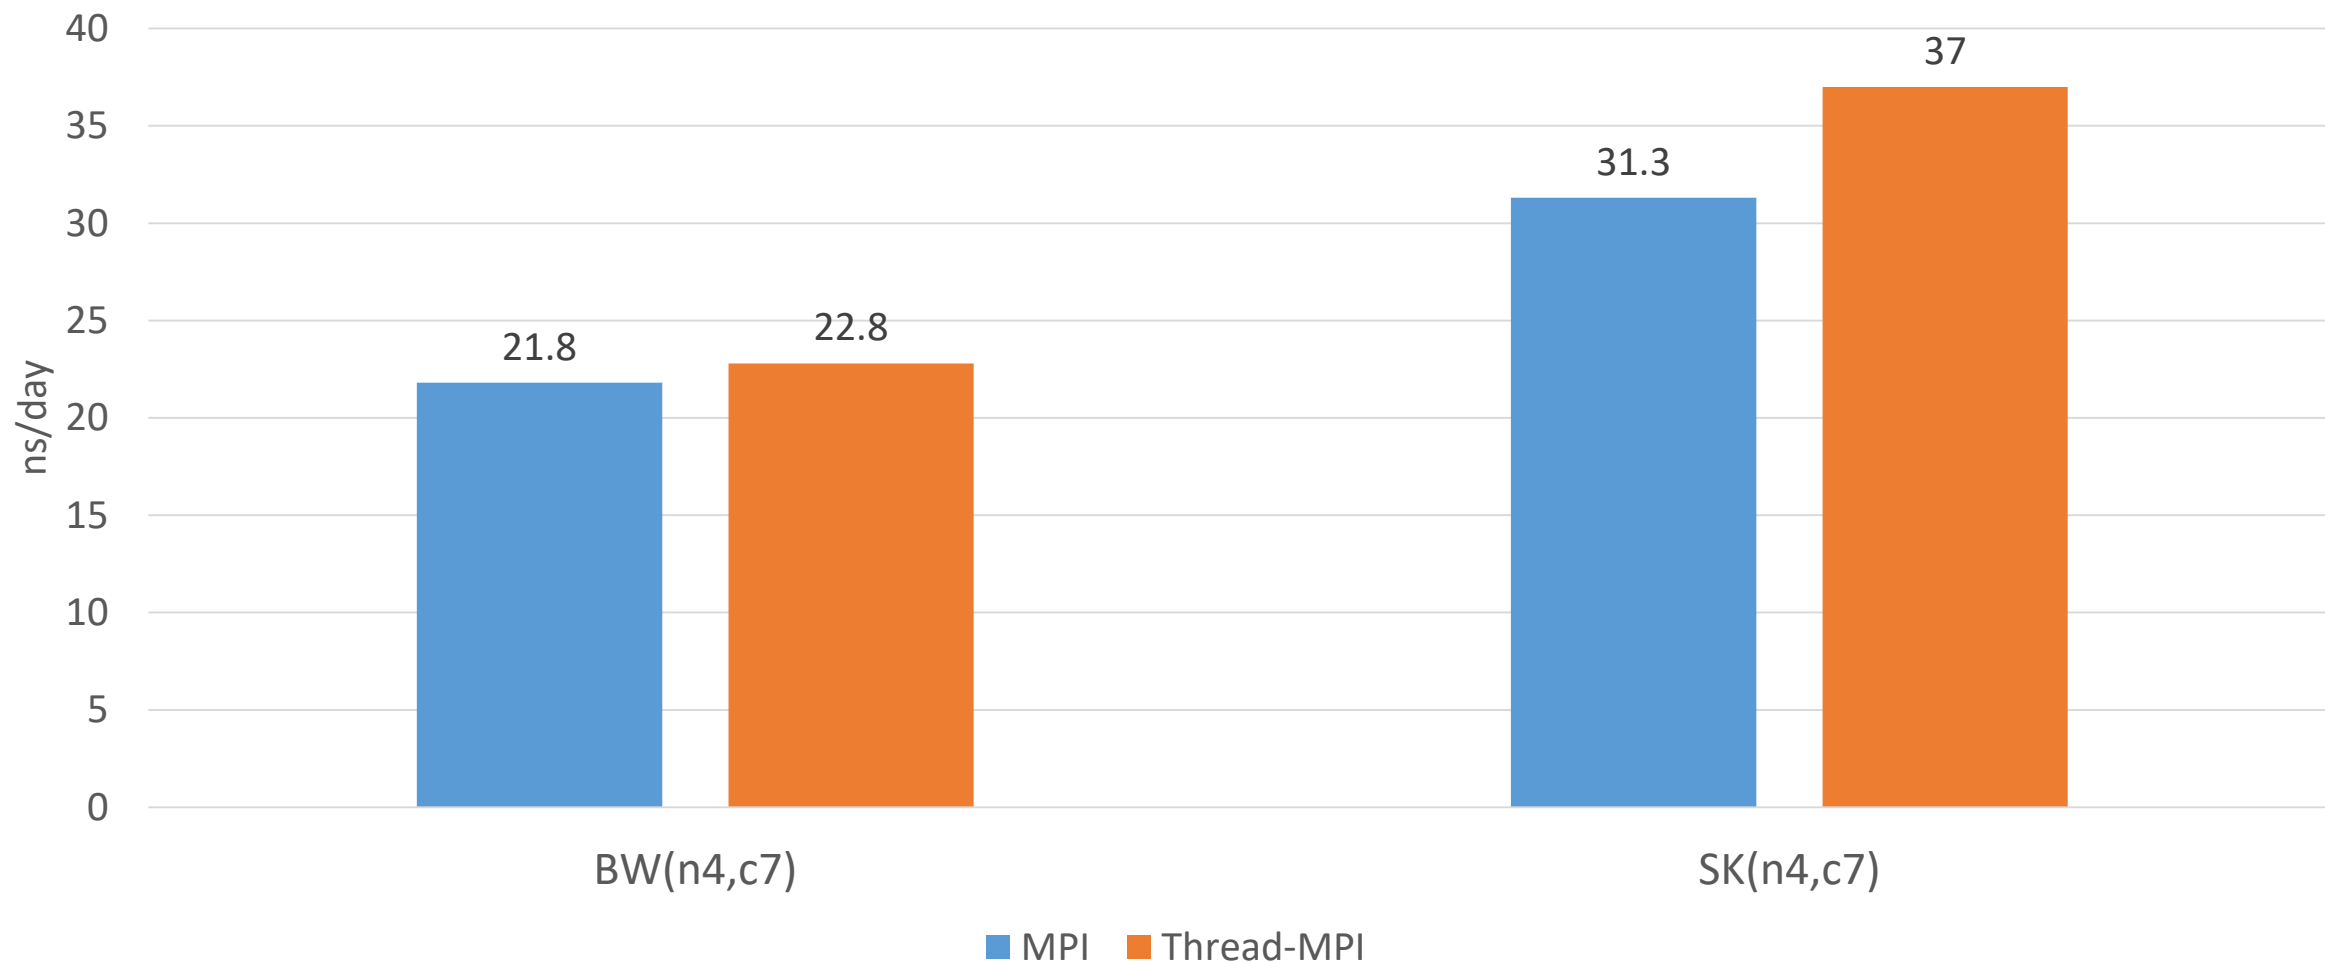
GROMACS (single node)

GROMACS-GPU



GROMACS (single node)

GROMACS/2021



GROMACS (single node-offload)

- (-n4,-c7) `srun gmx_mpi mdrun -ntomp 7 -dlb yes` will use 4 GPUs:

Using 4 MPI processes

Using 7 OpenMP threads per MPI process

On host b-cn1309.hpc2n.umu.se 4 GPUs auto-selected for this run.

Mapping of GPU IDs to the 4 GPU tasks in the 4 ranks on this node:

PP:0,PP:1,PP:2,PP:3

- (-n4,-c7) `gmx mdrun -gputasks 0123 -nb gpu -pme gpu -npme 1 -ntmpi 4 -dlb yes` will use 4 GPUs:

npme must be 1

Using 4 MPI threads

Using 7 OpenMP threads per tMPI thread

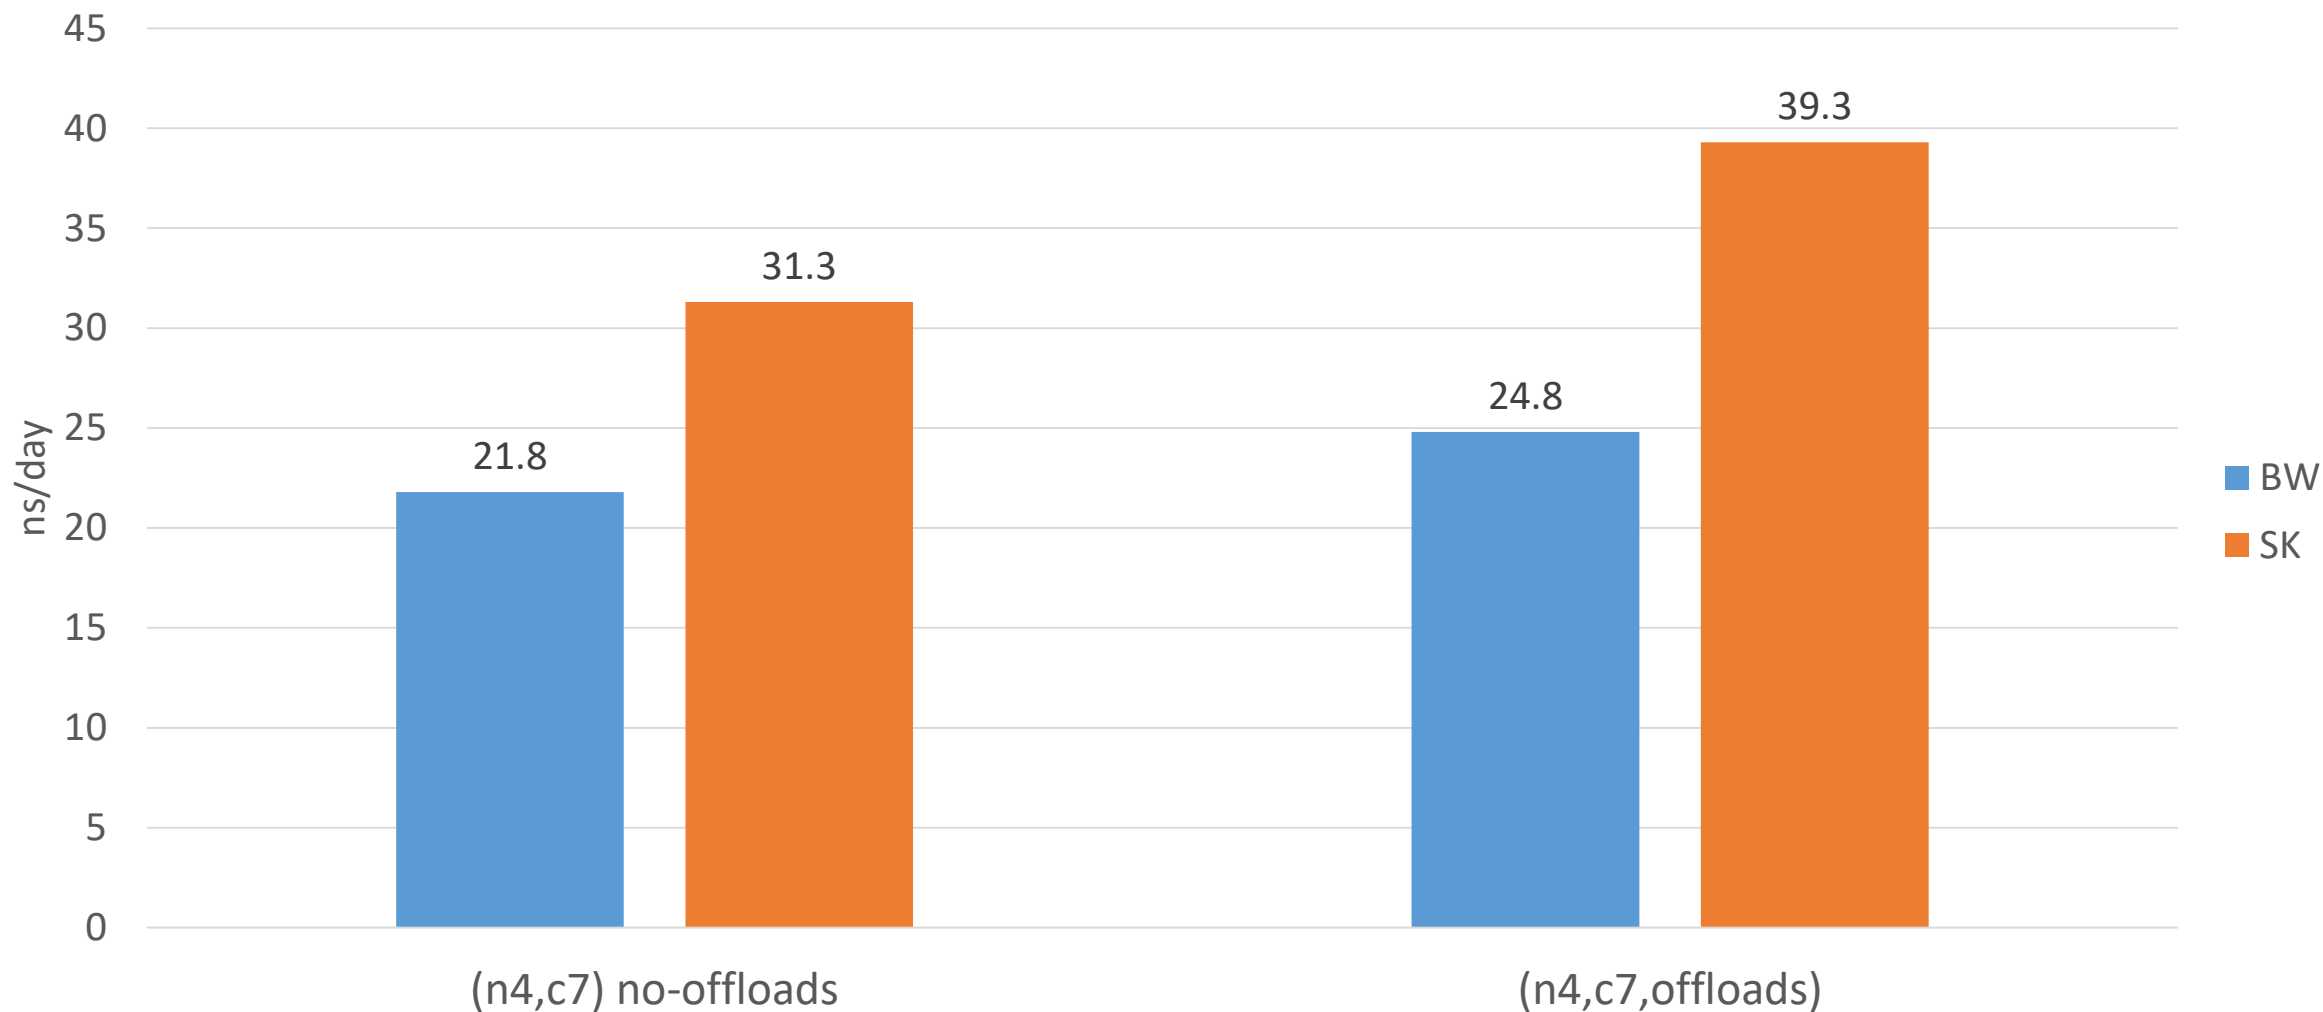
On host b-cn1102.hpc2n.umu.se 4 GPUs user-selected for this run.

Mapping of GPU IDs to the 4 GPU tasks in the 4 ranks on this node:

PP:0,PP:1,PP:2,PME:3

GROMACS (single node-offload)

GROMACS/2021 (MPI)



GROMACS (multi node-offload)

- (-N2,-n8,-c7) `srun gmx_mpi mdrun -nb gpu -pme gpu -npme 1 -ntomp 7 -dlb yes` will use 4 GPUs:

Using 7 OpenMP threads per MPI process

On host b-cn1106.hpc2n.umu.se 4 GPUs user-selected for this run.

Mapping of GPU IDs to the 4 GPU tasks in the 4 ranks on this node:

PP:0,PP:1,PP:2,PP:3

PP tasks will do (non-perturbed) short-ranged interactions on the GPU


PME tasks will do all aspects on the GPU

- There are some limitations regarding the tasks that can be offloaded, please check:

<http://manual.gromacs.org/current/user-guide/mdrun-performance.html>

GROMACS (multi node-offload)

GROMACS/2019



GROMACS

- PLUMED, threaded MPI version is not supported.

Instead of:

```
gmx mdrun -ntmpi X
```

Use:

```
mpirun gmx_mpi ...
```

GROMACS

- Additional information:

Recently, GROMACS 2021 added new features for offloading more tasks to the GPU, for instance bonded interactions (**-bonded gpu**) and updates (**-update gpu**). One can also move hard parallelizable PME parts to cpu (**-pmefft cpu**)

Resources

Manual for the current version:

<https://manual.gromacs.org/current/user-guide/mdrun-performance.html>

ENCCS information

<https://enccs.github.io/gromacs-gpu-performance/md-algorithm/>

GROMACS

Resources

GROMACS tutorials:

<http://www.mdtutorials.com/gmx/>

BIOEXCEL

<https://bioexcel.eu/>

LAMMPS (MPI)

```
#!/bin/bash
#SBATCH -A staff
#Asking for 10 min.
#SBATCH -t 02:10:00
#Number of nodes
#SBATCH -N 1
#Ask for 28 processes
#SBATCH -n 28
#SBATCH --exclusive

#Load modules necessary for running LAMMPS
ml GCC/8.3.0 OpenMPI/3.1.4
ml LAMMPS/3Mar2020-Python-3.7.4-kokkos

#Execute LAMMPS
srun Imp -in step4.1_equilibration.inp
```

LAMMPS (OpenMP)

```
#!/bin/bash
#SBATCH -A staff
#Asking for 10 min.
#SBATCH -t 02:10:00
#Number of nodes
#SBATCH -N 1
#Ask for 28 processes
#SBATCH -n 14
#SBATCH -c 2
#SBATCH --exclusive

#Load modules necessary for running LAMMPS
ml GCC/8.3.0 OpenMPI/3.1.4
ml LAMMPS/3Mar2020-Python-3.7.4-kokkos
export OMP_NUM_THREADS=2
#Execute LAMMPS
srun lmp -in step4.1_equilibration.inp
```

~3x slower than MPI only version

LAMMPS


Resources

LAMMPS tutorials

<https://lammps.sandia.gov/tutorials.html>

LAMMPS - OpenMM - CHARMM

LAMMPS - OpenMM - CHARMM



Scratch directory

- If the program writes frequently data, you can get an additional speed up by using the local scratch directory:

```
parent=/proj/nobackup/projdir/somedir  
rsync -avh $parent /scratch
```

```
cd /scratch/somedir
```

```
namd2 +p 28 +setcpuaffinity input.inp > output.dat
```

```
rsync -avh /scratch/somedir/. $parent/.
```

```
rm -rf /scratch/somedir
```

~10-15% speedup

Kebnekaise, standard nodes: **171 GB**
Kebnekaise, GPU nodes: **171 GB**
Kebnekaise, Largemem nodes: **352 GB**
(a few of them have 391 GB)

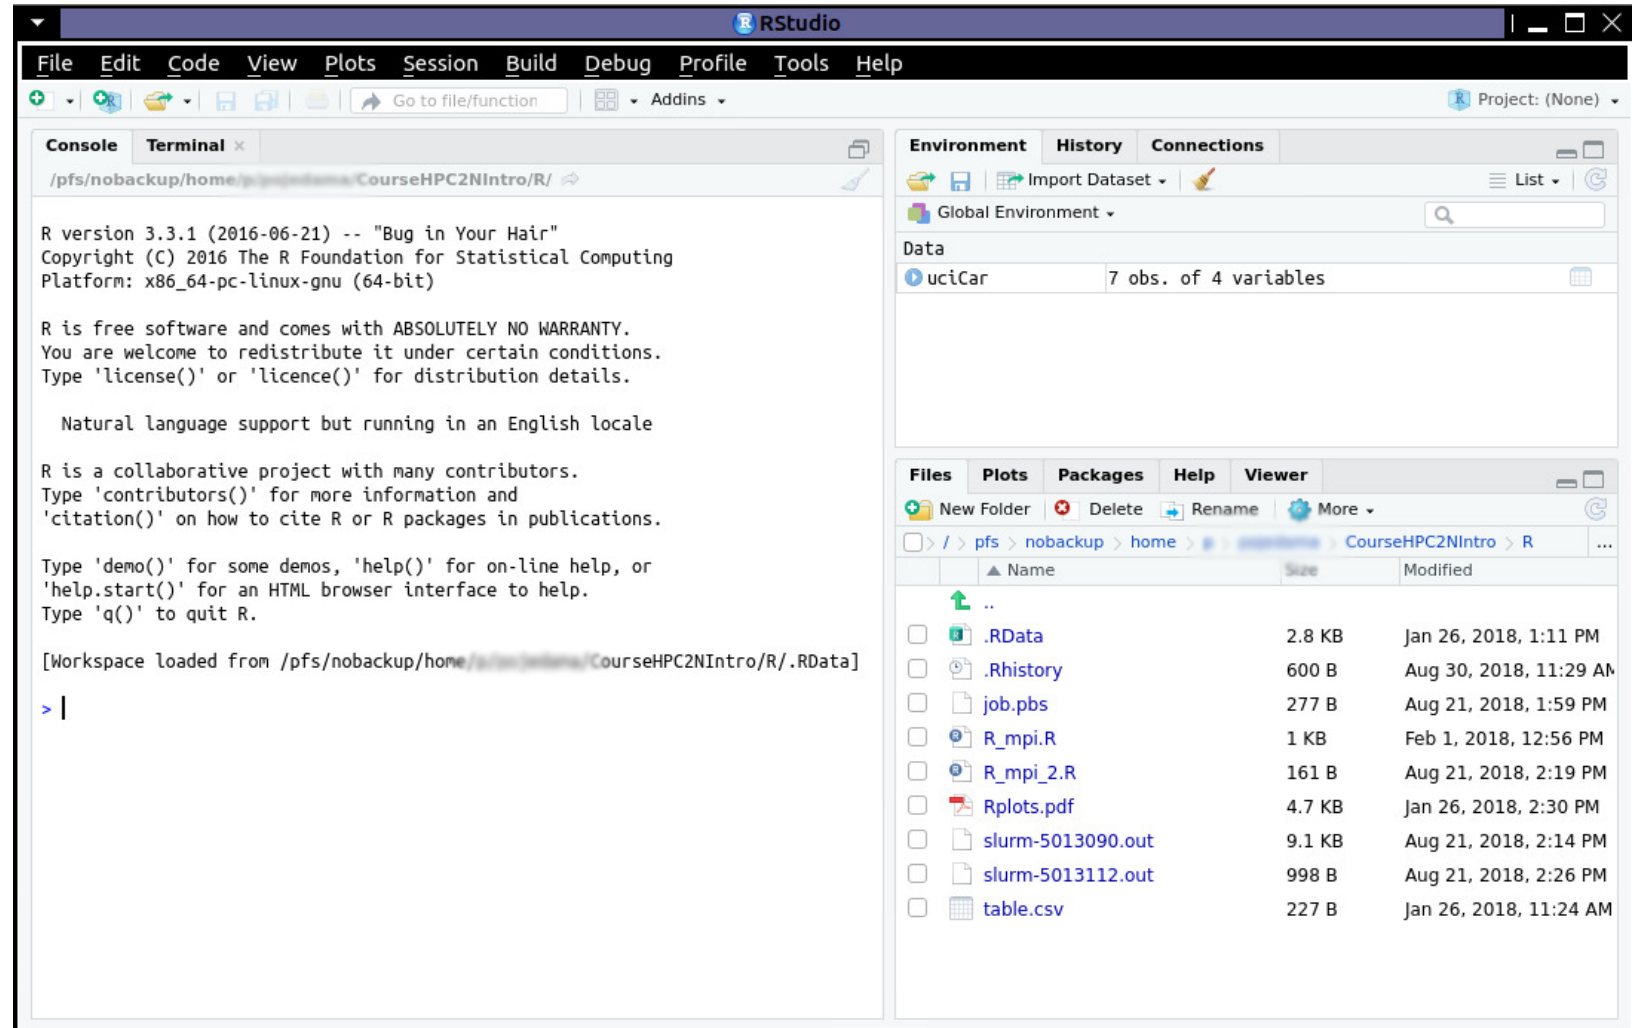
Plugins/Tools

Note: not all the tools referenced here are installed on our systems. It is indicated with the labels:
I=installed, MI=manually installed

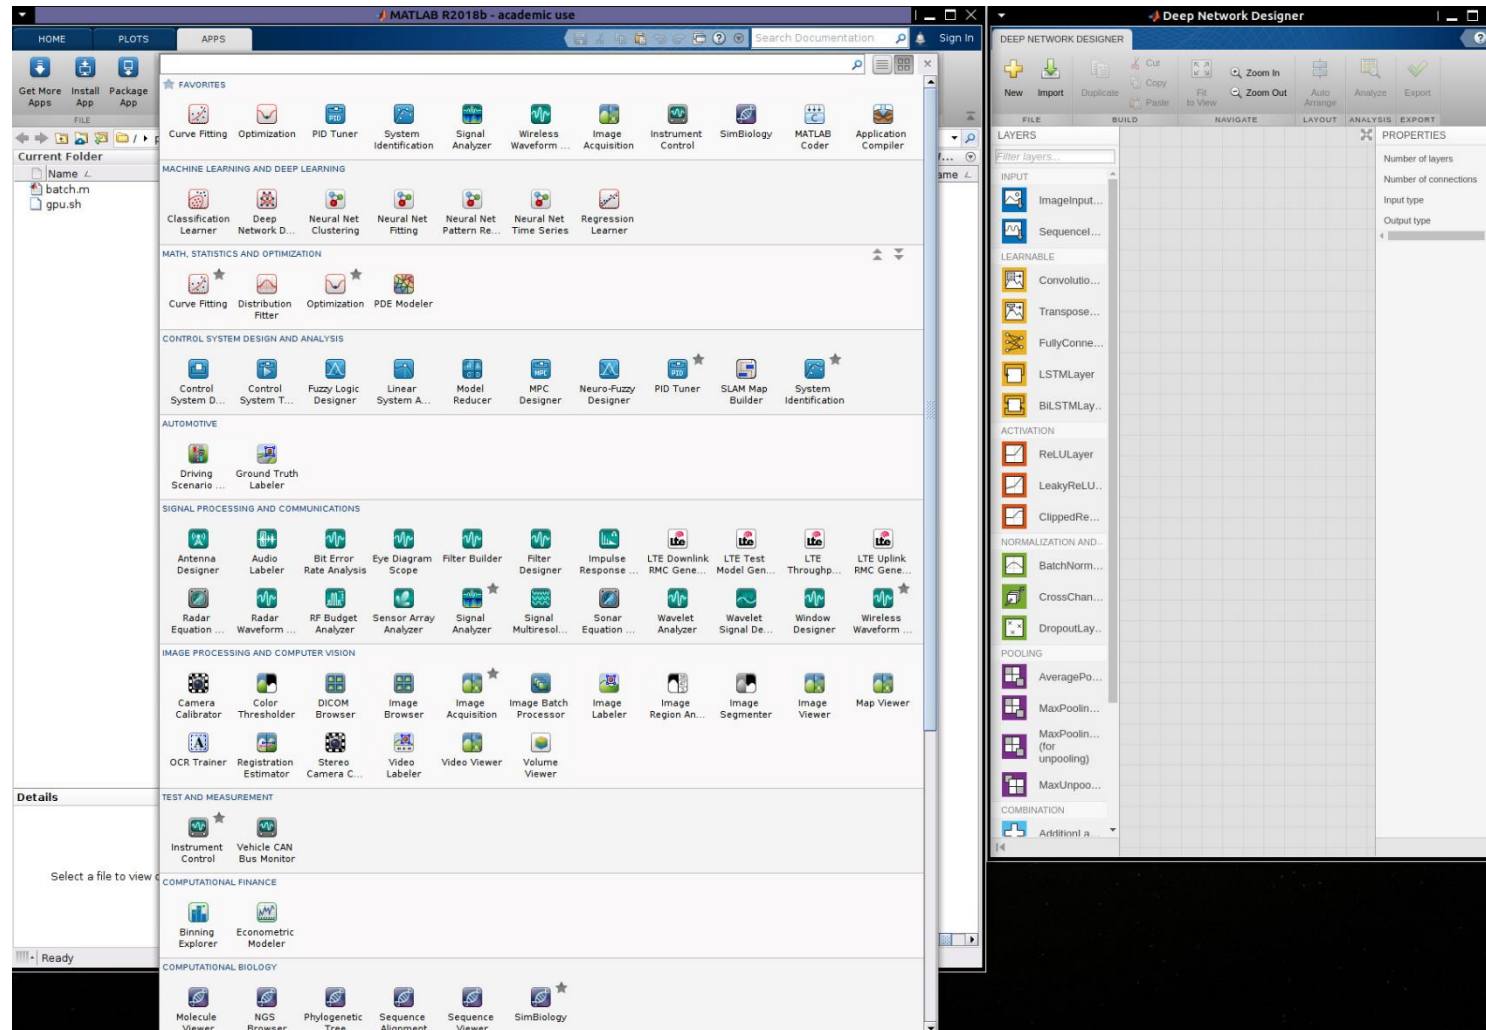
R/Rstudio (I)

```
$ml icc/2017.1.132-GCC-6.3.0-2.27  
impi/2017.1.132 ifort/2017.1.132-  
GCC-6.3.0-2.27
```

```
$ml R/3.3.1
$rstudio
```



MATLAB (I)



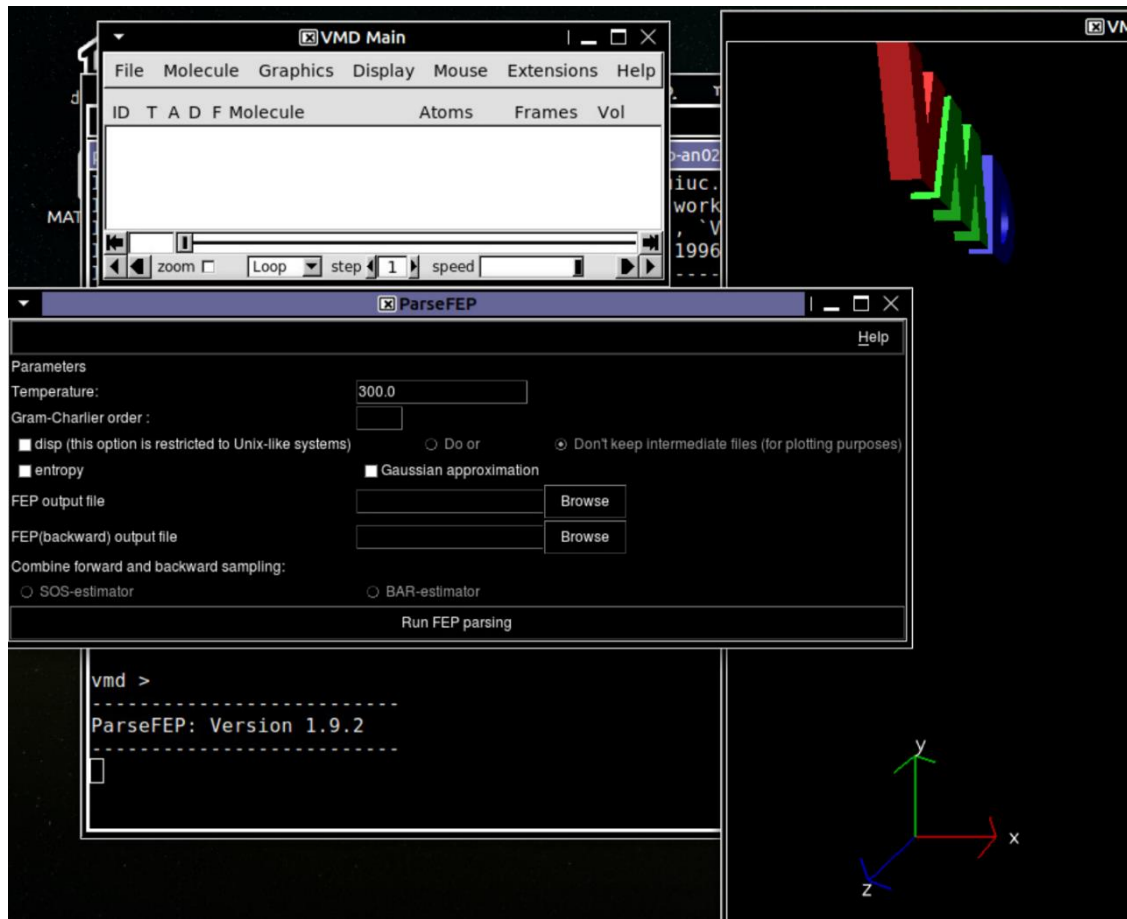
\$mI MATLAB/2018b

Includes improved tools
for Machine Learning and
Deep Learning

PRL, 120, 143001 (2018)

Further information: <https://www.hpc2n.umu.se/resources/software/matlab>

Alchemical simulations (I)




parseFEP in VMD software

$$H(\lambda) = \lambda H_B + (1 - \lambda) H_A$$

$$\Delta G = G_B - G_A = \sum_{\lambda=0}^1 -RT \ln \left\langle e^{-\Delta H'/RT} \right\rangle_{\lambda}$$

Chem. Rev., 93, 2395-1417 (1993)



Dual topology. Figure taken from:

<http://www.ks.uiuc.edu/Training/Tutorials/namd/FEP/tutorial-FEP.pdf>

NEMD-2D (I)

Simulation performed using LAMMPS

```
pair_style lj/cut ${rc}

# shear rate defined relative to perpendicular dimension


variable xyrate equal ${srate}/ly

fix 1 all nvt/sllod temp $t $t 0.1
fix 2 all deform 1 xy erate ${xyrate} remap v
```

More information:

<https://github.com/lammps/lammps/tree/master/examples/VI>
SCOSITY

Similar NEMD simulations can be performed with GROMACS:
Mol. Sim., 36, 560-567 (2010).



QM/MM simulations

- CHARMM SQM (MI)

!Topology and coordinates files

!Define the QM region

define subs sele segid mnp show end

!Specify the QM method


mndo remo sele subs show end -

GLNK sele none show end amd char -3 -

switched DXLBond NORD 9 NSTEpSCF 7

!Perform the dynamics

dynamics cpt leap restart time 0.0005 nstep 40000 ...



Adapted from Methods in Mol. Biology, vol. 924.
G. Groenhof.

QM/MM simulations

- CHARMM SQM (MI)

!Topology and coordinates files

!Define the QM region

define subs sele segid mnp show end

!Specify the QM method

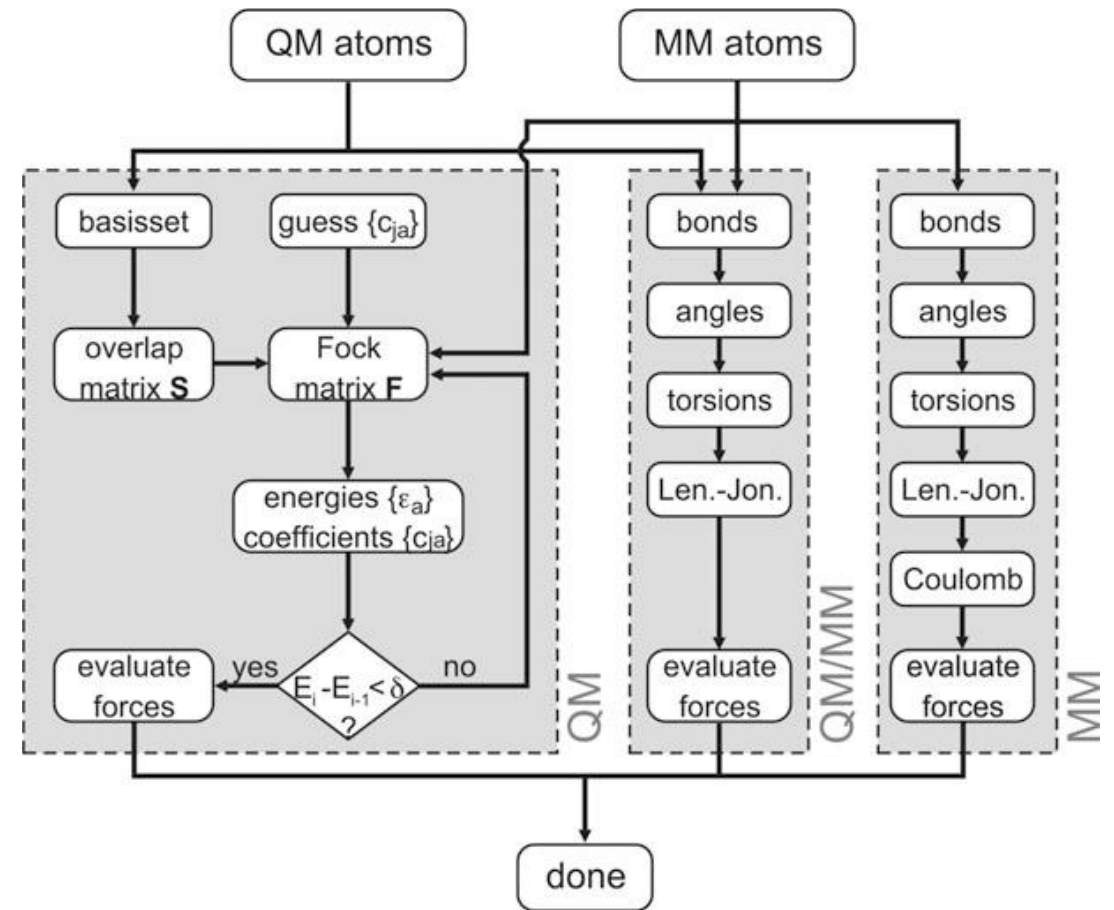
mndo remo sele subs show end -

GLNK sele none show end amd char -3 -

switched DXLBomd NORD 9 NSTEpSCF 7

!Perform the dynamics

dynamics cpt leap restart time 0.0005 nstep 40000 ...

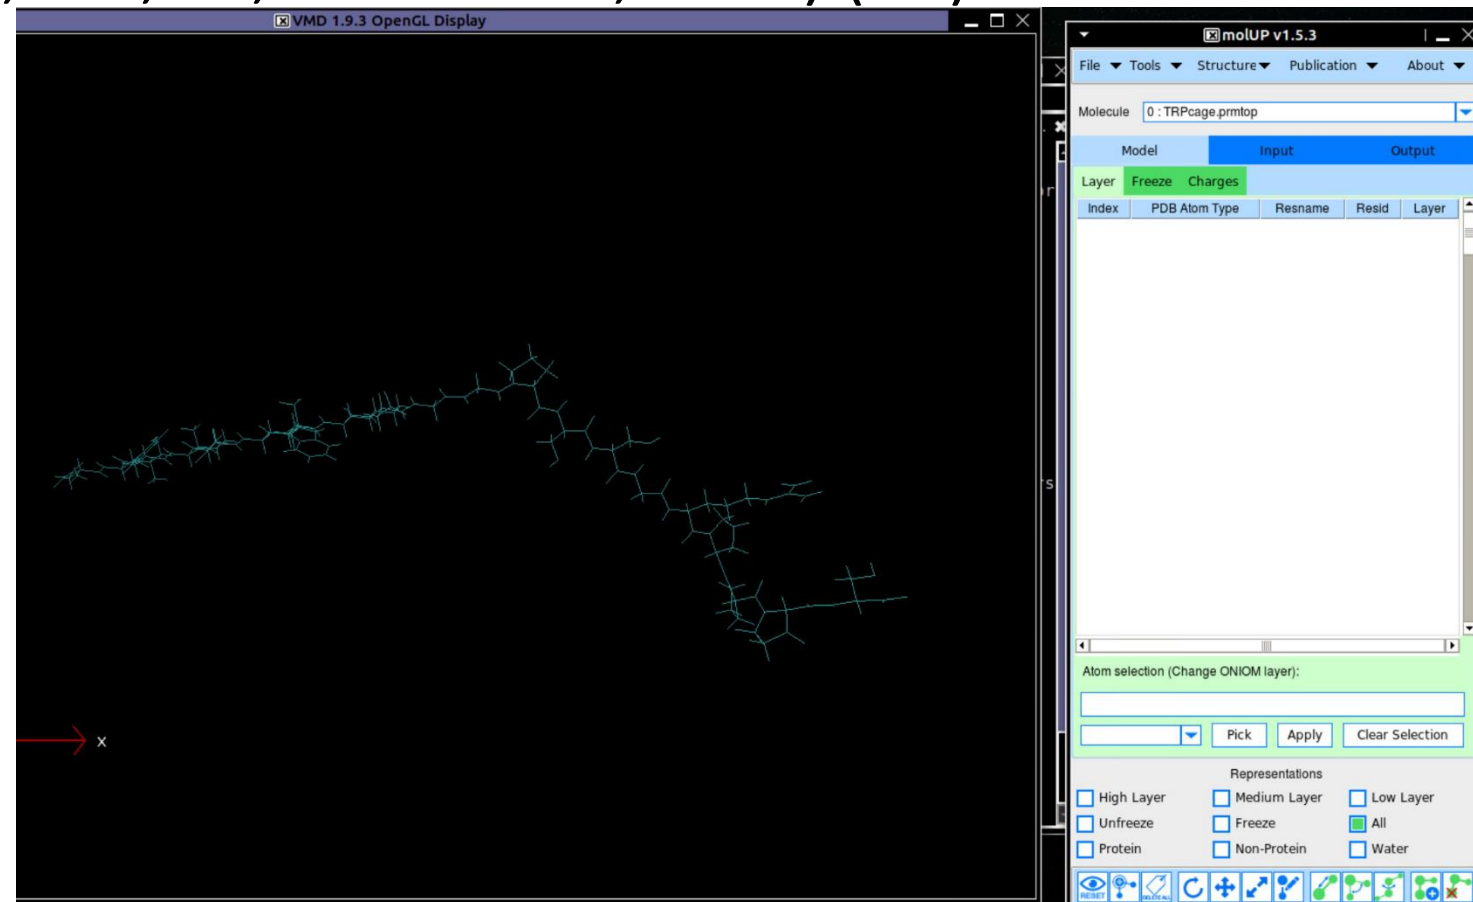


Adapted from Methods in Mol. Biology, vol. 924. G. Groenhof.


QM/MM simulations

- CHARMM SQM
- molUP, (Amber, JCC, 39, 1344-1353, 2018) (MI)

molUP plugin for QM/MM
simulations in VMD



Constant pH simulation (I)



NAMD config file taken from:

<http://www.ks.uiuc.edu/Training/Tutorials/>

Method described in: JCTC, 13, 5933-5944 (2017)

<http://ambermd.org/tutorials/>

```
source ../../lib/namdcph/namdcph.tcl
```

...

```
parameters $toppar_dir/par_all36_solvent.prm
# Load constant-pH specific topology files
cphConfigFile $toppar_dir/conf_cph36_prot.json
topology      $toppar_dir/top_cph36_prot.rtf
```

```
# We will be running multiple pH values sorted into their
# own directories, but
# otherwise using the same naming scheme.
```

```
set pH 3.5
```

```
pH $pH
```

```
outputname    $pH/ace_asp_nme_prod0
```

```
cphMDBasename $pH/namdcph.md
```

```
cphSwitchBasename $pH/namdcph.sw
```

```
# With the current settings this implies 1 ps between
# switching attempts,
# which will be 15 ps in length.
```

```
#
```

```
cphNumstepsPerSwitch 7500
```

```
cphRun 500 5
```

Constant pH simulation (I)

- Recent progress in this field, Cruzeiro et. al, JCP, 149, 072338 (2018), *Redox potential replica exchange molecular dynamics at constant pH in AMBER: Implementation and validation*

Analysis

- StreaMD: Advanced analysis of molecular dynamics using R, JCC, 39, 1666 (2018). (MI)
- <http://thegrantlab.org/bio3d/tutorials/normal-mode-analysis> (MI)
- <http://www.ks.uiuc.edu/Training/Tutorials/>
- https://github.com/rjdkmr/do_x3dna DNA structural analysis (GROMACS) (MI)

Validity test in MD

RESEARCH ARTICLE

Testing for physical validity in molecular simulations

Pascal T. Merz, Michael R. Shirts*

PLOS ONE (2018)

Suggestions

- Check if the software version you are trying is MPI (-n), OpenMP (-c), or hybrid (-n -c). Also, if you are using a GPU version.
- Is the cutoff radius for long-range interactions appropriate? PME?
- Could KNL nodes speed up your workflow?
- Keep your .bashrc file clean
- Upon asking a question to the support team: Make a folder with the case which displays the issue. Including all relevant files would be useful.
- Don't forget to cite HPC2N and SNIC when publishing.
- Monitor your simulation on the fly with the command: "job-usage Job_ID"

Terminology

- BW: Broadwell nodes
- SK: Skylake nodes
- KNL: Knights Landing nodes
- MTS: Multiple Time Step algorithm